



A modern database interface for C++

Erik Smith

WHY DO WE CARE

some data points

- sqlite installed base > # smart phones
- JDBC (java's database API) was introduced in 1998

PREVIOUS PROPOSAL WORK

- N3886 - Johann Anhofer - A Proposal to add a Database Access Layer to the Standard Library
- N3612 - Thomas Neumann - Desiderata of a C++11 Database Interface
- N3458 - Thomas Neumann - Simple Database Integration in C++
- N3415 - Bill Seymour - A Database Access Library

THE RELATIONAL MODEL

"rowset"

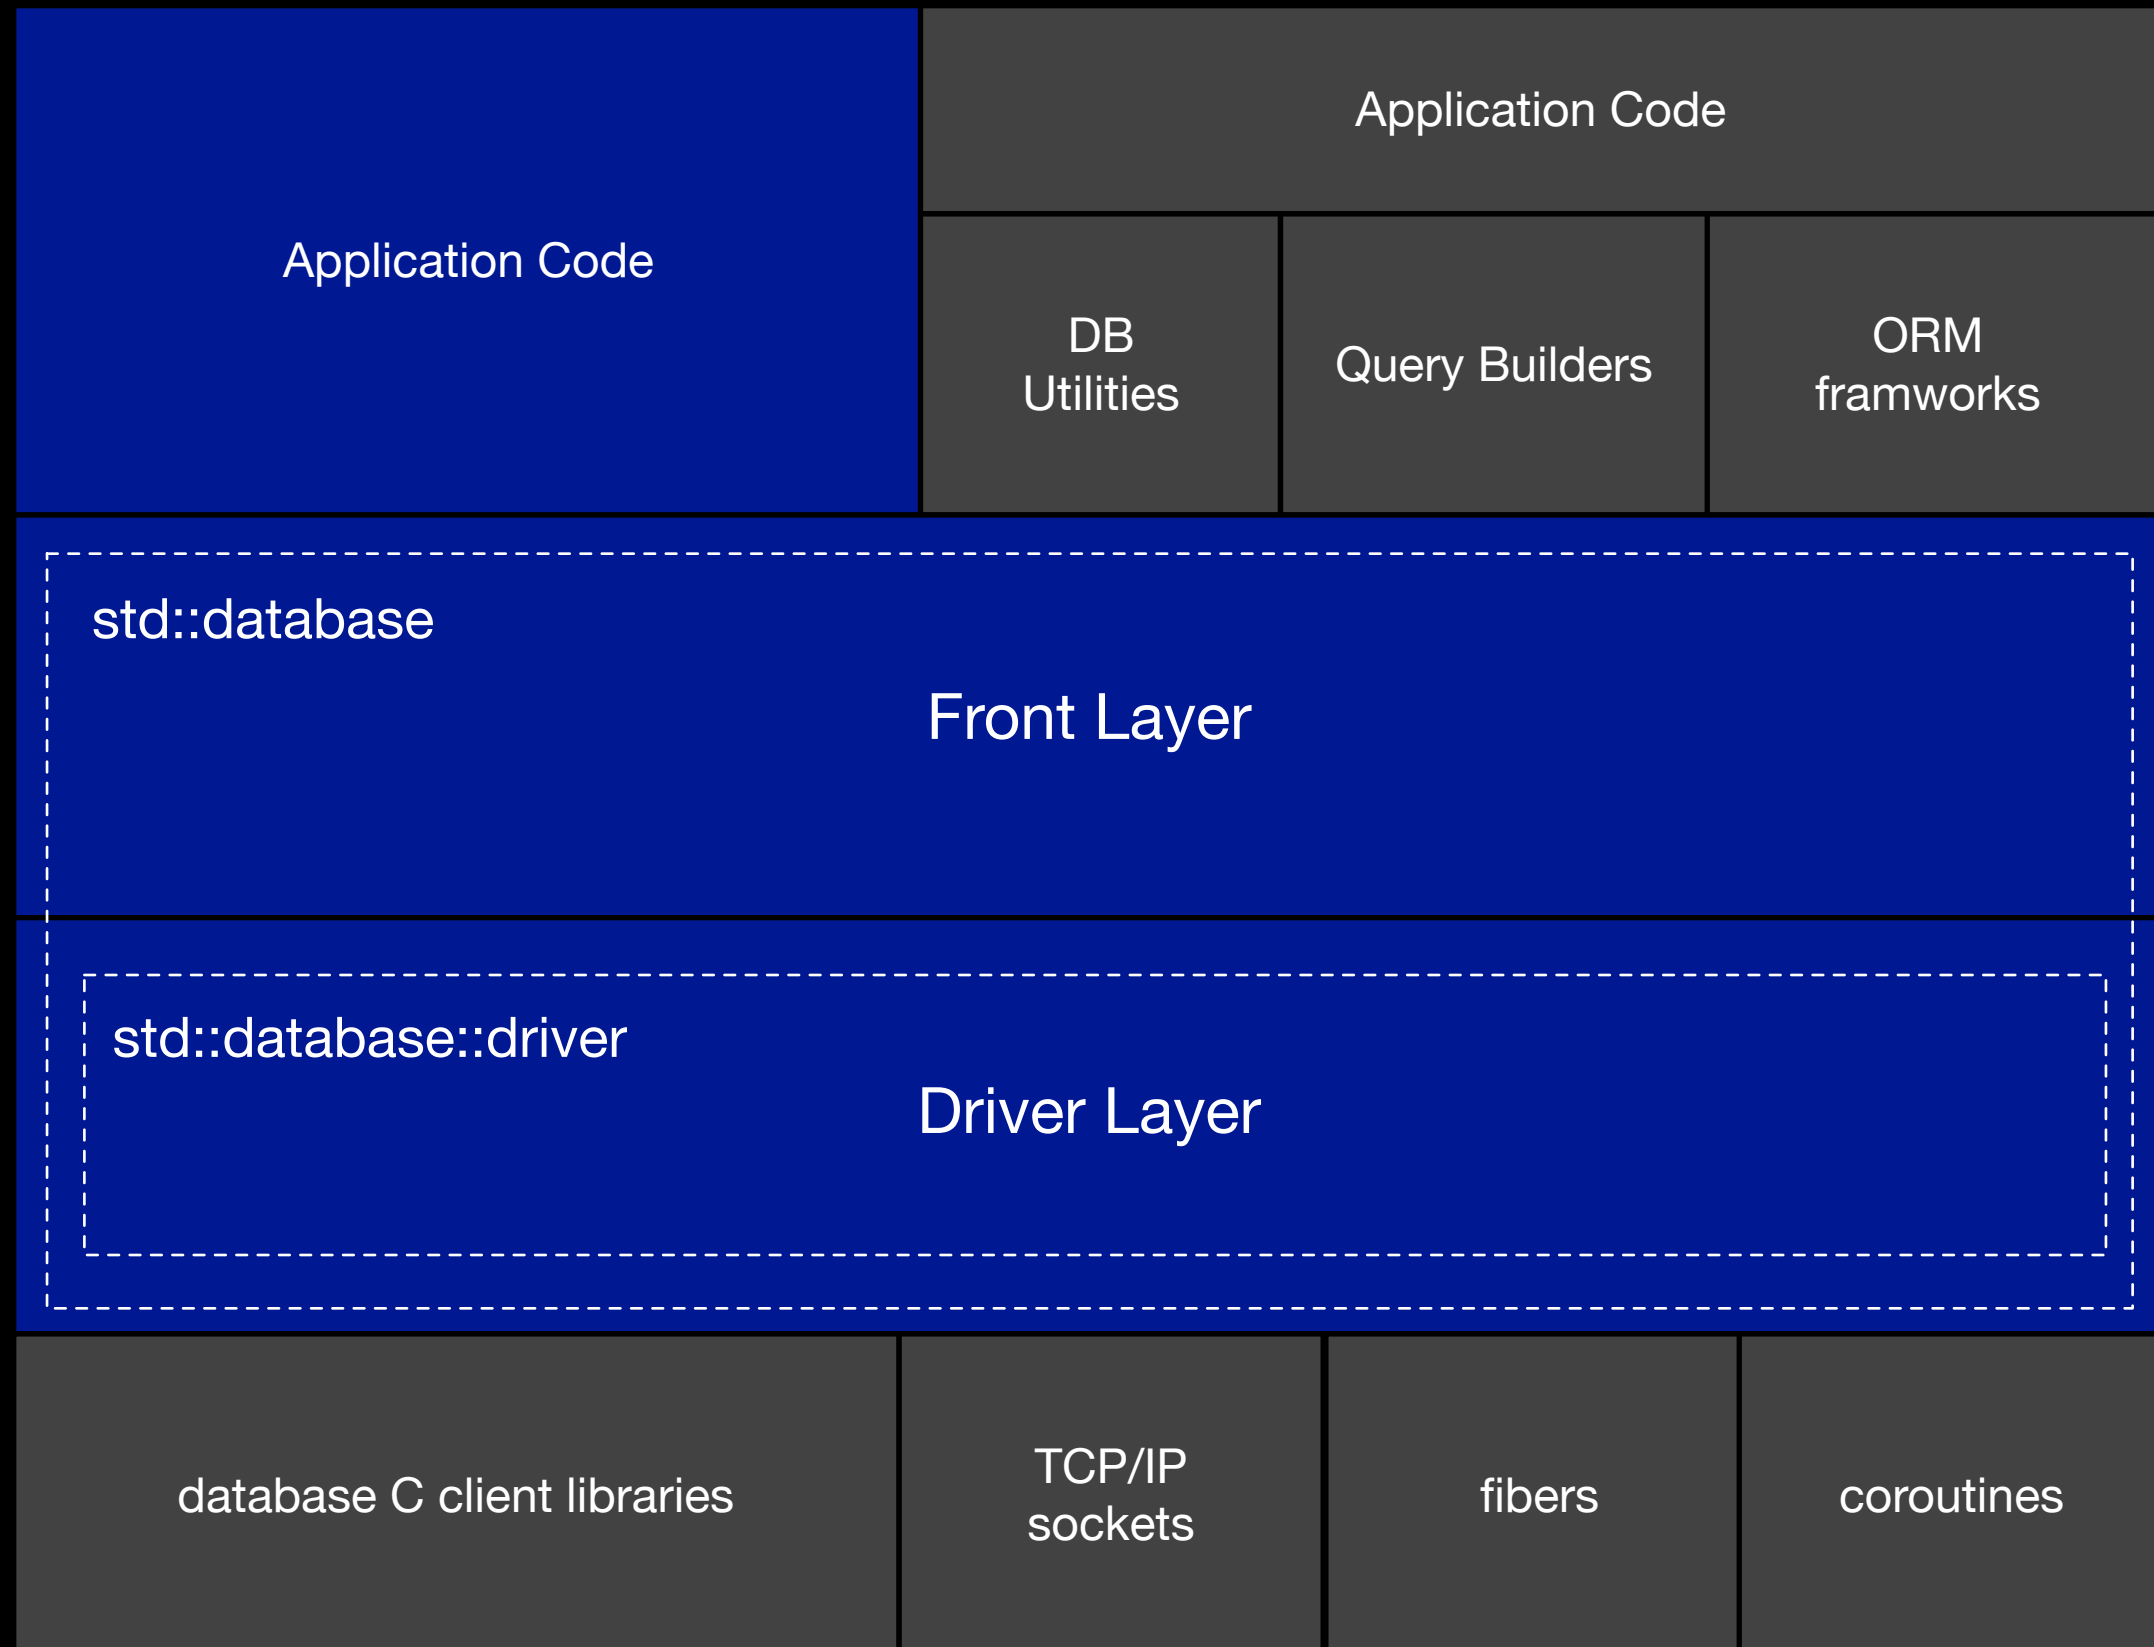
Table

columns	PERSON			
	ID INT	LAST_NAME VARCHAR(30)	FIRST_NAME VARCHAR(30)	BIRTHDATE DATE
	1	Yokomoto	Akiko	1990-05-03
	2	Green	Marjorie	1972-02-06
rows	3	Hoffman	Paul	1995-07-01

SQL query

```
select last_name from person
where birthdate < "1991-01-01"
```

THE CLIENT STACK



EXAMPLE

```
using cppstd::mysql;
auto db = create_database("mysql://server/db");
db
    .query("select * from person")
    .rows()
    .write(cout);
```

no explicit types

INCLUDE & NAMESPACE ALIAS

details for particular implementation

```
#include <cppstd/db/mysql/database.h>
```

```
using cppstd::db::mysql;
```

```
auto db = create_database("mysql://server/db");
```

```
db
```

```
    .query("select * from person")
```

```
    .rows()
```

```
    .write(cout);
```

COMPOSABLE

uses a classic method chaining approach

```
class connection {  
    auto query(const string& q) {  
        ...  
        return *this;  
    };  
};
```

recent Uniform Function Call Syntax (UCFS) proposals
(N4474) support this approach

```
con.query(q) failover to query(con,q);
```


DIRECT VS POLY

direct: using one specific database type
(no indirection cost)

```
#include <cppstd/db/mysql/database.h>  
using cppstd::mysql;  
auto db = create_database("mysql://server/db");
```

poly: choose any registered database type at run-time
(more on poly later)

THE TYPES

database	top level context (shared)
connection	connection to server (per thread)
statement	query execution, prepared statements, input binding
columnset	meta-data for columns
column	meta-data for specific column
rowset	accessor for query results
row	row accessor (proxy) type
field	field accessor (proxy), type conversion.

EXAMPLE IN EXPANDED ("CLASSIC") FORM

```
using namespace cppstd::mysql;
auto db = create_database("mysql://server/db");
auto con = db.connection();
auto stmt = con.query("select * from person");
auto rows = stmt.rows();
```

```
for (auto i = rows.begin(); i != rows.end(); ++i) {
    for(int c = 0; c != row.width(); ++c) {
        auto field = row[c];
        cout << "value: " << field << "\n";
    }
    cout << "\n";
}
```

DUAL METHOD STYLES

by type name

db

.connection(uri)

.statement()

.rowset()

natural style

db

.connect(uri)

.query(sql)

.rows()

keeping both for now as they both make sense

FLEXIBLE

these are equivalent

```
db.connection().statement("select * from t").execute().rows()
```

```
db.connection().query("select * from t").execute().rows()
```

```
db.connection().query("select * from t").rows()
```

```
db.query("select * from t").rows()
```

hard to get wrong

WHY REFERENCE TYPES?

a non-nested scope example (think MVC)

```
auto get_customers_by_zip(const string& zip) {  
    auto db = create_database("mysql://server/db");  
    db  
        .query("select from person where name=?", zip)  
        .rows();  
}  
  
auto local_customers = get_customers_by_zip("92122");  
    rowset is safe to escape  
    rowset holds shared state in std::shared_ptr
```

DATABASE

```
auto db = create_database;
```

```
auto db = create_database("mysql://server/db");           // default URI
```

```
auto con = database.connection();
```

```
// uses default URI
```

```
auto con = database.connection("mysql://server2/db");
```

```
// specific URI
```

```
auto con = database.connection("server");
```

```
// named source
```

```
auto con = database.create_connection();
```

```
// new connection
```

```
auto con = database.create_connection("mysql://server2/db");
```

```
db.query("insert into person values(1, 'joe');");
```

connection: returns same connection per thread

create_connection: new connection for same thread

CONNECTION STRING

URI based string

```
db.connection("mysql://server/db?username=app")
```

Named source

```
db.connection("mysql")
```

Custom Resolver 

config file

```
"databases": [  
  {  
    "name": "mysql",  
    "type": "mysql",  
    "server": "127.0.0.1",  
    "database": "test",  
    "username": "",  
    "password": ""  
  },  
]
```


CONNECTION

```
auto db = connection.database();      // parent accessor
```

```
connection.auto_commit(false);        // for transactions
```

```
connection.begin();
```

```
connection.save();
```

```
connection.commit();
```

```
connection.rollback();
```

```
connection.isolationLevel();
```

TRANSACTIONS & SCOPE EXIT

```
auto con1 = db.connection("server1").auto_commit(false);  
auto con2 = db.connection("server2").auto_commit(false);
```

```
auto goodbye = make_scope_exit([]() {    // proposal N4189  
    con1.rollback();  
    con2.rollback();  
});
```

```
con1.begin.query("insert into account(id,amount) values(1,-500000)");  
con2.begin.query("insert into account(id,amount) values(2, 500000)");
```

```
con1.commit();  
con2.commit();
```

STATEMENT

statement facilitates prepared statements / input binding

```
auto stmt = con
```

```
.query("insert into person(id, name) values(?,?)");
```

```
stmt.execute(1, "Doug");
```

```
stmt.execute(2, "Cathy");
```

```
stmt.execute(3, "Robert");
```

essential for security (SQL injection attacks)

ROWSET

accessors

<code>row.width()</code>	<code>// number of columns in row set</code>
<code>row.columns()</code>	<code>// column metadata</code>
<code>row.length()</code>	<code>// number of results (at end or materialized)</code>

what about row traversal?

APPROACHES TO ROWSET TRAVERSAL / ACCESS

- `next()`, `done()`
- `front()`, `pop_front()`, `empty()`
- iterators?

ITERATORS!

rowset uses InputIterators

forward only, only current result accessible

RANGE BASED FOR

```
auto db = create_database("mysql://server/db");  
auto rows = db.query("select * from person").rows()
```

```
for (auto row : rows) {  
    for (col : row.columns()) {  
        cout << "value: " << row[col] << "\n";  
    }  
    cout << "\n";  
}
```

RANGE BASED FOR

rows is an Iterable type with iterator proxies - refine

for (auto row : rows)	// works - copy ok!
for (auto&& row : rows)	// also fine
for (auto& row : rows)	// error as expected
for (const auto& row : rows)	// error as expected
for (row : rows)	// N3853, not today

ITERATOR OVERHEAD

Are iterators for row traversal bad for performance?

```
class iterator {
```

```
...
```

```
    bool operator==(const rowset_iterator& rhs) const {
```

```
        return
```

```
            (rowset_ && !rowset_->done()) ==
```

```
            (rhs.rowset_ && !rhs.rowset_->done());
```

```
    }
```

```
...
```

```
};
```

Not that bad for InputIterators, but
still have to account for 4 cases at
runtime

STD LIBRARY ALGORITHM EXAMPLE 1

```
auto db = create_database();
auto rng = db.query("select name,score from score").rows();

auto key = "carol";
auto i = find_if(rng.begin(), rng.end(), [key](auto row) {
    return row["name"] == key; });

if (i != rng.end()) {
    cout << "found row: " << *i << "\n";
}
```

STD LIBRARY ALGORITHM EXAMPLE 2

```
auto db = create_database();  
auto rng = db.query("select name,score from score").rows();  
  
int sum = accumulate(rng.begin(), rng.end(), 0, [](int sum, auto row) {  
    return sum + row["score"].as<int>();});  
  
cout << "sum: " << sum << "\n"; // 150
```

ITERATOR ISSUES

- proxy iterators will be problematic for rowset types beyond InputIterator

RANGES !

Ranges are a proposed addition to the C++ standard library from Eric Niebler.

Ranges address existing issues with iterators, particularly with composability, and bring a whole new class of functionality with range adapters.

RANGE SUPPORT

A rowset is an InputRange

Iterable Concept

begin(s) and end(s) defined, not necessarily same type

Range Concept

Extends Iterable

lightweight

default constructible, copyable, assignable, destructible

InputRange

iterator type is InputIterator

- No pure Iterable type exposed in interface
- A rowset is a range
- Element data is accessed through internal shared_ptr
- Only the current row is accessible.

RANGE OPTIMIZATIONS

standard library (C++98) input iterator comparison:

```
bool operator==(const rowset_iterator& rhs) const {  
    return  
        (rowset_ && !rowset_->done()) ==  
        (rhs.rowset_ && !rhs.rowset_->done());  
}
```

comparison optimized for ranges:

```
bool operator==(const sentinel& rhs) const {  
    return !rowset_ || rowset_->done();  
}
```

RANGES: PASS RANGE INSTEAD OF PAIR

```
auto db = create_database();
auto rng = db.query("select name,score from score").rows();

auto key = "carol";
auto i = find_if(
    rng,
    [key](auto row) {return row["name"] == key; });

if (i != rng.end()) {
    cout << "found row: " << *i << "\n";
}
```

Compose the range in the argument to find_if

RANGE ADAPTORS

`view::filter` like a lazy version of `std::copy_if`

```
auto db = create_database();  
int min_score = 20;  
auto rng = view::filter(  
    db.query("select name,score from score").rows(),  
    [min_score](auto row) {return row["score"] > min_score; });
```

`view::filter` returns a range that is set up to filter lazily
great for proxy InputRanges like rowset
`std::copy_if` won't work here

ROW

```
auto field = row[0];           // by column index
```

```
auto field = row["FIRST_NAME"]; // by column name
```

```
auto field = row["first_name"]; // case insensitive
```

```
auto field = row[column];      // by column
```

FIELD

```
cout << field << "\n";           // stream inserter=  
field.as<T>()                      // type T  
field.as<int>()                    // int  
field.as<long>()                   // long  
field.as<std::string>()            // std::string  
field.as<std::string_view>()       // std::string_view  
  
field.variant()                    // as std::variant (nothrow)  
field.optional<T>()                // as std::optional<T> (nothrow)
```

FIELD ACCESSORS

<u>field.is_null()</u>	// is the value null
field.name()	// name of column
field.type()	// type enumeration

ROW FIELD INDEXING

```
auto field = row["name"];    // easy but less efficient
```

```
auto field = row[1];         // efficient but less readable
```

MORE EFFICIENT ROW FIELD INDEXING

more efficient, readable?

```
auto id_idx = rows.columns()["id"];  
auto name_idx = rows.columns()["last_name"];
```

```
for(auto row : rows) {  
    auto id = row[id_idx];  
    auto name = row[name_idx];  
}
```

harder to maintain

SINGLE ROW QUERIES WITH into

```
string a;
```

```
int b;
```

```
date d;
```

```
variant d;
```

```
db
```

```
    .query("select a,b,c,d from table")
```

```
    .into(a,b,c,d);
```

ROW LEVEL into

better living though variadics

```
auto rows =  
    .query("select id, name from person")  
    .rows();
```

```
int id;
```

```
string name;
```

```
for (auto row : rows) row.into(id, name);
```


INTO RANGE

```
variant rng;  
query(sql).into(rng);  
query(sql).row().into(rng);
```

need to disambiguate rowset vs row

SERIALIZATION

equivalent syntax to for object serialization

```
person p
```

```
db.query(sql).into(p);
```

```
rows.into(p)
```

```
row.into(p)
```

```
row >> p;
```

Important use case for reflection in the standard

HANDLE ACCESSORS

```
auto con = db.connection();
```

```
auto mysql = connection.handle();           // typed as MYSQL*
```

```
auto mysql = db.connection.handle();         // lifetime fail
```



OUTPUT BINDING

ID	NAME
2	JOE

- A C level buffer interface
- RowSet handles internally

ARRAY OUTPUT BINDING

```
auto rows = con
    .query("select * from t1")
    .rows();

int sum;
for(auto r : rows)
    sum += r[0].as<int> + r[1].as<int>;
```

1000 row table

A INT	B INT
0	1
1	2
2	3
...	...
999	1000

203 ms

ARRAY OUTPUT BINDING

```
auto rows = con
    .row_array_size(100)
    .query("select * from t1")
    .rows();

int sum;
for(auto r : rows)
    sum += r[0].as<int> + r[1].as<int>;
```

1000 row table

A INT	B INT
0	1
1	2
2	3
...	...
999	1000

32 ms

ARRAY OUTPUT BINDING

```
auto rs = con
    .row_array_size(500)
    .query("select * from t1")
    .rows();

int sum;
for(auto r : rows)
    sum += r[0].as<int> + r[1].as<int>;
```

1000 row table

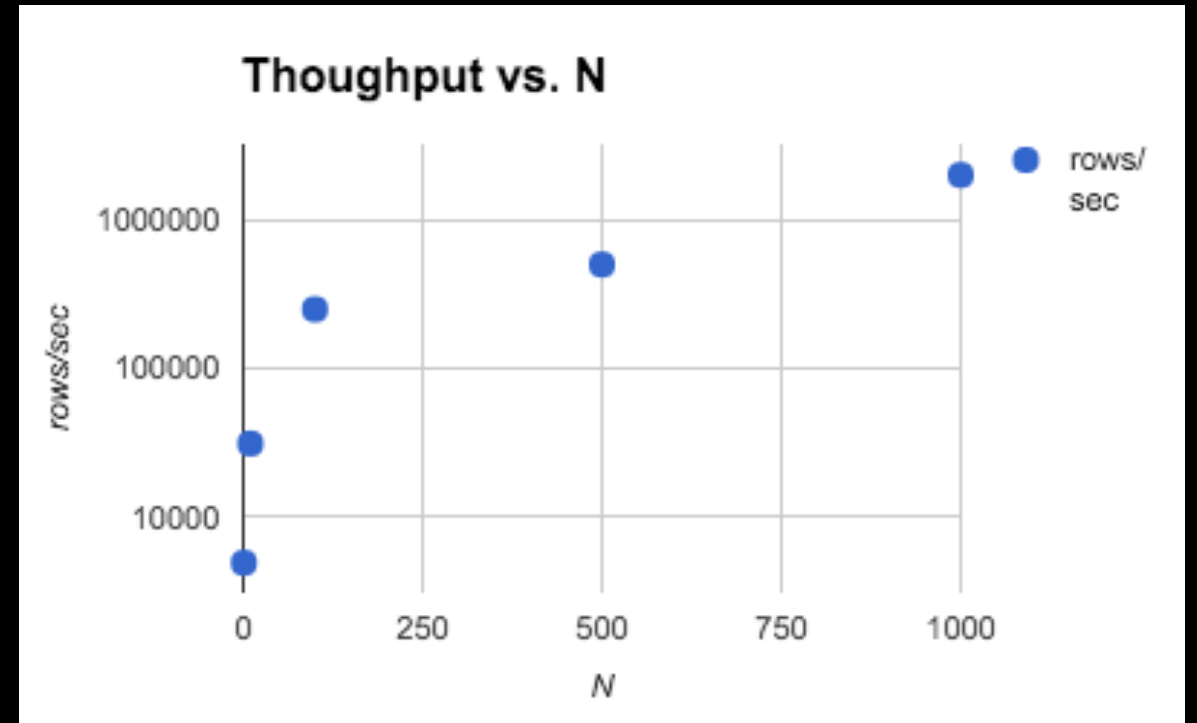
A INT	B INT
0	1
1	2
2	3
...	...
999	1000

2 ms

ARRAY OUTPUT BINDING

```
auto rs = con
    .row_array_size(1000)
    .query("select * from t1")
    .rows();
```

```
int sum;
for(auto r : rows)
    sum += r[0].as<int> + r[1].as<int>;
```



601 μ s

2M rows/sec

400X faster

CONTIGUOUS ARRAY BINDING

some databases can bind column arrays into the same contiguous memory (struct binding / skip parameters)

DETACHED ROWSETS

```
auto rows = detached_rows();
```

- detachable rowset can detach from connection
- rowset is now a RandomAccessRange
- no additional copying
- rowset caching enabler

INPUT BINDING

same as before

```
auto stmt =  
    create_database("mysql://server/db")  
    .query("insert into table(id, name) values(?,?)");  
  
for(auto d : data) stmt.execute(d.id, d.name);
```

ARRAY INPUT BINDING

```
auto stmt =  
    create_database("mysql://server/db")  
    .query("insert into table(id, name) values(?,?)")  
    .row_array_size(1000);  
  
for(auto d : data) stmt.query(d.id, d.name);
```

huge performance win

TYPE CONVERSION

- Two layers (driver & front end)
- Native driver binding is default

POLICIES

```
struct my_policy {...}  
auto db = database<my_policy>;
```


- Custom allocators
- Pluggable connection pool
- assert on release build for cross/illegal type conversions
- Scoped types (no RC)
- Omit handle accessors

UTILITY EXAMPLE: JOIN

```
auto inventory = db.create_connection().query("select id,* inventory").rows();  
auto orders = db.create_connection().query("select * from orders order by id").rows();  
auto parts = db.create_connection().query("select * from parts order by id").rows();
```

```
auto joined_rows = ordered_natural_join(inventory, orders, parts);
```

```
for(auto row : joined_rows) {  
    int id, order_id, part_id;  
    inventory.into(id,...);  
    orders.into(order_id,...);  
    parts.into(part_id,...);  
}
```



overlapping key condition: all tables ordered consistently
takes $O(1)$ space

an approach to the "multiple hierarchies" problem

Use IOT/clustered indexes when tables are big

JOIN ALGORITHMS

// ordered join with specific column conditions:

```
auto j = ordered_join(inventory.join(orders).on(orders, "part_id");
```

other more general join possibilities with detached
(RandomAccess) rowsets

POLYMORPHIC INTERFACE

Direct Interface

```
auto db = create_database();  
auto con = db.connection("mysql://server/db");
```

Poly Interface

```
auto db = cppstd::poly::create_database();  
auto con1 = db.connection("mysql://server/db");  
auto con2 = db.connection("sqlite://file.sqlite");
```

poly could conceivably be in the standard library

POLY: ADDING DRIVERS

```
#include <cppstdadb/poly/database.h>
```

```
using cppstdadb::poly;
```

```
database.register<cppstdadb::sqlite::database>();
```

```
database.register<cppstdadb::mysql::database>();
```

```
database.register<cppstdadb::oracle::database>();
```

```
database.register<cppstdadb::postgres::database>();
```

REFERENCE IMPLEMENTATION TEST SUITE

```
test_all<cppstd::mysql::database>("mysql");  
test_all<cppstd::sqlite::database>("sqlite");  
test_all<cppstd::oracle::database>("oracle");
```

...

- Templated test framework
- Runs test twice: once direct, ones through poly driver
- Runs carefully in sandbox database

IMPLEMENTATION DETAILS

two layer design

Front End

- Handles reference counting details for all types
- Defines all interface functions
- Consolidates calls to the driver
- Manages state
- Connection pooling

Driver

- Implement driver specific details

DRIVER INTERFACE

```
namespace cppstdb { namespace mysql {
```

```
using database =
```

```
    std::database::basic_database<  
        driver<default_policy>,default_policy>;
```

```
auto create_database() {return database<default_policy>();}
```

```
template<class policy> class driver {
```

```
    class database {...}
```

```
    class connection {...}
```

```
    class statement {...}
```

```
    class result {...}
```

```
}
```

```
}}
```

Type name correspondence
between front and
driver layers

DRIVER INTERFACE

```
template<class policy> class driver {  
    class database {...};  
    class connection {  
        connection(database *db, source src, allocator *a) {...}  
    };  
    class statement {  
        statement(connection *con, string sql, allocator *a) {...}  
    };  
    class result {  
        result(statement *stmt, allocator *a) {...}  
    };  
}
```

CHALLENGE: POLY EXECUTE

```
template<typename... Args> auto execute(Args... args) {  
    // unpack variants  
    // form 2nd variadic template call in driver layer  
}
```

#1 Avoid requiring the drivers to handle Variant arguments

#2 Need to transport arguments from one variadic template query call to another at *run time*

Run time  Compile Time

APPROACH

Front End

```
stmt.execute("joe", date(2015,2,1), 42)
```



vector<variant_t>

V<string>

V<int>

V<date>



variadic_dispatch (from unpack_variants call with 3 args)



Driver

```
unpack_variants(v[0],v[1],v[2])
```



```
stmt.execute("joe", Date(2015,2,1), 42)
```

string

date

int

Is this a good idea?




2 minutes to compile

REALITY CHECK

10 types

5 arguments

$P(10,5) = 30240$ function calls



Alternative:

```
auto execute(std::vector<variant_t>& args); // an additional call
```

OR

execute must support variants

not a big issue

NON-BLOCKING INTERFACE

Can be implemented for underlying database clients that
that support NBIO



"We're Gonna Need A Bigger Database"

REFERENCE IMPLEMENTATION ROADMAP

current support (limited depth)



phase II (soon)



phase III



QUESTIONS

github: <https://github.com/cruisercoder/cppstdadb>

twitter: @cruisercoder