

C++ Costless Abstractions

the Compiler View

Proudly made in *Namek* by serge-sans-paille

/me

Serge « sans paille » Guelton

```
$ whoami  
sguelton
```

- R&D engineer at QuarksLab on compilation for security
- Associate researcher at Télécom Bretagne

He that started it all

«C++ Is my favorite garbage collected language because it generates so little garbage» – Bjarne Stroustrup

Zero-Cost Abstraction

«C++ enables zero-overhead abstraction to get us away from the hardware without adding cost.» – Bjarn Stroustrup

FACT: There's always a cost.

But the compiler (and the language) pay it in **complexity**.
You don't pay it in **execution time**.

constexpr \neq costless

```
constexpr int fibo(int v) {  
    return v < 2 ? v : (fibo(v-1) + fibo(v-2));  
}  
int main() {  
    static_assert(fibo(26) == 121393, "ok");  
    // std::cout << fibo(26) << std::endl;  
    return 0;  
}
```

- Compilation times with constexpr: ~0.9s
- Compilation times without constexpr: ~0.3s

constexpr trivia

Changing Fibo number from 26 to 27

```
% time clang++ constexpr.cpp -std=c++14
constexpr.cpp:7:17: error: static_assert expression is not an integral constant expression
    static_assert(fibo(27) == 196418, "ok");
                   ^~~~~~
constexpr.cpp:2:27: note: constexpr evaluation hit maximum step limit; please use a smaller constant
constexpr int fibo(int v) {
```



```
}
```

A Brief Word about Clang+LLVM

- Modern compiler infrastructure
- Historical focus on C-like languages
- Uses a typed Internal Representation

Used to illustrate all the following examples

Disclaimer

The results in these slides are relative to the compiler used (clang++-3.6), the OS (Linux) and the arch (amd64).

Almost nothing is guaranteed by the standard

Functions

Functions are an essential piece of abstraction:

- Give a name to a block of code
- Avoid redundancy
- Abstract with respect to types → overloading

Functions

But a function call implies a performance penalty

- **Saving context**
- **Jumping (twice)**

And they are **over-used** in the STL:

- `std::copy`
- calls `std::__copy_move_a2`
- that calls `__copy_move_a`
- that calls `__copy_move`
- that may calls `__builtin_memmove`

Inlining - Input

```
struct foobar {  
    int doit(int a) const { return a+1;}  
};  
  
inline int foo(int a) {  
    return a+1;  
}  
int bar(foobar const &fb, int val) {  
    return foo(fb.doit(val));  
}
```

Inlining - Output

```
define i32 @_Z3barRK6foobari(%struct.foobar* nocapture readnone dereferen
    %1 = add nsw i32 %val, 2
    ret i32 %1
}
```

Inlining

Q: is `inline` useful?

A: not much more than an hint (+ODR)

Q: how to take control?

A: `-mllvm -inline-threshold=225` or `-inlinehint-threshold=42`

Q: how to really take control?

A: `__attribute__((always_inline))`

Value semantic

```
void foo(double & f) {  
    f += 1.;  
}  
double bar(double f) {  
    return f + 1.;  
}
```

Value semantic

```
define void @_Z3fooRd(double* nocapture dereferenceable(8) %f) {  
    %1 = load double* %f  
    %2 = fadd double %1, 1.000000e+00  
    store double %2, double* %f  
    ret void  
}  
  
define double @_Z3bard(double %f) {  
    %1 = fadd double %f, 1.000000e+00  
    ret double %1  
}
```

Const Ref or Value?

```
static int foo(int const& v) __attribute__((noinline)) { return v; }  
static int bar(int v) __attribute__((noinline)) { return v; }  
  
int caller(int v) {  
    return foo(v) + bar(v);  
}
```


Const Ref or Value?

Compiled with -O2

```
define internal fastcc i32 @_ZL3fooRKi(i32* nocapture readonly dereferenc
    %1 = load i32* %v
    ret i32 %1
}

define internal fastcc i32 @_ZL3bari(i32 %v) {
    ret i32 %v
}
```

Const Ref or Value?

Compiled with -O3

```
define internal fastcc i32 @_ZL3fooRKi(i32 %v.val) {  
    ret i32 %v.val  
}  
  
define internal fastcc i32 @_ZL3bari(i32 %v) {  
    ret i32 %v  
}
```

struct const &

```
struct pack {  
    unsigned a,b,c;  
};  
  
static int foo(pack const& v) __attribute__((noinline)) { return (v.a + v.b); }  
int caller(pack const& v) {  
    return foo(v);  
}
```

Passing struct by value

Changing signature!

```
define internal fastcc i32 @_ZL3fooRK4pack(i32 %v.0.0.val, i32 %v.0.1.val, i32 %v.0.2.val) {
    %1 = add i32 %v.0.1.val, %v.0.0.val
    %2 = add i32 %1, %v.0.2.val
    %3 = udiv i32 %2, 3
    ret i32 %3
}
```

Passing struct by value

Tail call + unboxing

```
_Z6callerRK4pack:
    movl    (%rdi), %eax
    movl    4(%rdi), %esi
    movl    8(%rdi), %edx
    movl    %eax, %edi
    jmp     _ZL3fooRK4pack          # TAILCALL

_ZL3fooRK4pack:
    addl    %esi, %edi
    leal    (%rdi,%rdx), %ecx
    movl    $2863311531, %eax
    imulq   %rcx, %rax
    shrq    $33, %rax
    retq
```

Tag dispatching

```
struct tag0 {};  
struct tag1 {};  
static int foo(int v, tag0) __attribute__((noinline)) { return v + 0; }  
static int foo(int v, tag1) __attribute__((noinline)) { return v * 2; }  
  
int caller(int v) {  
    return foo(v, tag0{}) + foo(v, tag1{});  
}
```

Tag dispatching

No useless argument

```
define internal fastcc i32 @_ZL3fooi4tag0(i32 %v) {  
    ret i32 %v  
}  
  
define internal fastcc i32 @_ZL3fooi4tag1(i32 %v) {  
    %1 = shl nsw i32 %v, 1  
    ret i32 %1  
}
```

Lambda

```
int (*foo)(int, int) = [](int x, int y) { return x + y; };  
int bar(int x, int y) {  
    return x + y;  
}
```


Lambda

Not different from a regular function?

```
define i32 @_Z3barii(i32 %x, i32 %y) {  
    %1 = add nsw i32 %y, %x  
    ret i32 %1  
}  
  
define internal i32 @"_ZN3$_08__invokeEii"(i32 %x, i32 %y) {  
    %1 = add nsw i32 %y, %x  
    ret i32 %1  
}
```

Lambda + capture

```
auto foo(int val) {  
    return [val](int x) { return x + val; };  
}  
auto bar() {  
    return [] (int x) { return x * 3 ; };  
}
```

Lambda + capture

Not different from its state!

```
define i32 @_Z3fooi(i32 %val) {  
    ret i32 %val  
}  
  
define void @_Z3barv() {  
    ret void  
}
```

struct, class = data

- Use `class` to **box** data with type information
- *Eventually* to associate treatment to data
- Manage the lifetime of data

Boxing values

```
struct A {  
    char a_ = 'a';  
};  
struct B : A {  
    char b_ = 'b';  
};  
B* foo() {  
    return new B();  
}
```

Boxing values

```
%struct.B = type { %struct.A, i8 }
%struct.A = type { i8 }

; Function Attrs: uwtable
define noalias %struct.B* @_Z3foov() {
    %1 = tail call noalias i8* @_Znwm(i64 2)
    %2 = bitcast i8* %1 to %struct.B*
    %3 = bitcast i8* %1 to i16*
    store i16 0, i16* %3
    store i8 97, i8* %1
    %4 = getelementptr inbounds i8* %1, i64 1
    store i8 98, i8* %4
    ret %struct.B* %2
}
```

Boxing values

```
_Z3foov:
    pushq    %rax
.Ltmp0:
    movl     $2, %edi
    callq    _Znwm
    movw     $0, (%rax)
    movw     $25185, (%rax)
    popq     %rdx
    retq
```

More Boxing/Unboxing

```
struct foo {  
    short val_;  
    struct {  
        short val_;  
        struct {  
            short val_;  
        } nest_;  
    } nest_;  
};  
  
foo bar(char val) {  
    foo a;  
    a.val_ = val;  
    a.nest_.val_ = val;  
    a.nest_.nest_.val_ = val;  
    return a;  
}
```


More Boxing/Unboxing

Fuse scalars!

```
define i48 @_Z3barc(i8 signext %val) {  
    %1 = sext i8 %val to i16  
    %2 = zext i16 %1 to i48  
    %3 = shl nuw i48 %2, 32  
    %4 = shl nuw nsw i48 %2, 16  
    %5 = or i48 %4, %2  
    %6 = or i48 %5, %3  
    ret i48 %6  
}
```

Member Functions

```
struct foo {  
    int bar() const;  
};  
  
int foo_bar(foo const& f) {  
    return f.bar();  
}
```

Member Functions

Just a function (very Pythonic!)

```
%struct.foo = type { i32 }

define i32 @_Z7foo_barRK3foo(%struct.foo* dereferenceable(4) %f) {
    %1 = tail call i32 @_ZNK3foo3barEv(%struct.foo* %f)
    ret i32 %1
}

declare i32 @_ZNK3foo3barEv(%struct.foo*)
```

Default Copy Constructor

```
struct foo {  
    int a,b,c,d,e,f,g,h,i,j,k;  
};  
  
foo a;  
  
void test(foo b) {  
    a = b;  
}
```

Default Copy constructor

Memcpy detected!

```
%struct.foo = type { i32, i32, i32, i32, i32, i32, i32, i32, i32, i32, i32 }

@a = global %struct.foo zeroinitializer

; Function Attrs: nounwind uwtable
define void @_Z4test3foo(%struct.foo* byval nocapture readonly %b) {
    %1 = bitcast %struct.foo* %b to i8*
    call void @llvm.memcpy.p0i8.p0i8.i64(i8* bitcast (%struct.foo* @a to i8*,
    ret void
}
```

Default Copy constructor

Memcpy? Too expansive!

```
struct foo {  
    int a,b,c,d;  
};  
  
foo a;  
  
void test(foo b) {  
    a = b;  
}
```

Default Copy constructor

```
define void @_Z4test3foo(i64 %b.coerce0, i64 %b.coerce1) {  
    store i64 %b.coerce0, i64* bitcast (%struct.foo* @a to i64*)  
    store i64 %b.coerce1, i64* bitcast (i32* getelementptr inbounds (%struct.foo, @a, #0), i64*)  
    ret void  
}
```

Copy Elision

```
struct S {  
    double x_, y_, z_;  
    S(double x) : x_{x}, y_{x}, z_{x} {}  
    S(S const&) = default;  
};  
  
S init() {  
    return S(1.);  
}  
  
S init2() {  
    S d(1.);  
    return d;  
}
```


CopyElision

For free thanks to the representation

```
define void @_Z4initv(%struct.S* noalias nocapture sret %agg.result) {  
    %1 = bitcast %struct.S* %agg.result to <2 x double>*  
    store <2 x double> <double 1.000000e+00, double 1.000000e+00> <2 x double> %1  
    %2 = getelementptr inbounds %struct.S, %agg.result, i64 0, i32 2  
    store double 1.000000e+00, double* %2  
    ret void  
}
```

CopyElision

Same code with a temporary

```
define void @_Z5init2v(%struct.S* noalias nocapture sret %agg.result) {  
    %1 = bitcast %struct.S* %agg.result to <2 x double>*  
    store <2 x double> <double 1.000000e+00, double 1.000000e+00> <2 x double>*, %1  
    %2 = getelementptr inbounds %struct.S, %struct.S*, @_Z5init2v, i64 0, i32 2  
    store double 1.000000e+00, double* %2  
    ret void  
}
```

Virtual method calls

```
struct Interface {  
    virtual int doit() = 0;  
    virtual ~Interface() {}  
};  
struct A : Interface {  
    int doit() final { return 0; }  
};  
  
int foo(Interface& a) {  
    return a.doit();  
}  
int bar(A& a) {  
    return a.doit();  
}
```

Virtual method calls

Here comes the **vtable**

```
define i32 @_Z3fooR9Interface(%struct.Interface* dereferenceable(8) %a)
    %1 = bitcast %struct.Interface* %a to i32 (%struct.Interface*)***
    %2 = load i32 (%struct.Interface*)*** %1
    %3 = load i32 (%struct.Interface**)** %2
    %4 = tail call i32 @%3(%struct.Interface* %a)
    ret i32 %4
}

define i32 @_Z3barR1A(%struct.A* nocapture readnone dereferenceable(8) %a)
    ret i32 0
}
```

Virtual method calls

Force virtual call

```
int foobar() {  
    A a;  
    Interface& b = a;  
    return b.doit();  
}
```

Virtual method calls

It's devirtualized!

```
define i32 @_Z6foobarv() {  
    ret i32 0  
}
```

Misc

Various bells and whistles in C++ you don't pay for

Initializer list

```
std::array<int, 16> foo = {0,1,2,3,4,5,6,7,8,9, 0xA, 0xB, 0xC, 0xD, 0xE,  
std::vector<int> bar = {0,1,2,3,4,5,6,7,8,9, 0xA, 0xB, 0xC, 0xD, 0xE, 0xF}
```


Initializer list

Static initialisation

```
@foo = global %"struct.std::array" { [16 x i32] [i32 0, i32 1, i32 2, i32
```

Initializer list

Runtime initialisation

```
@bar = global %"class.std::vector" zeroinitializer  
(...)
```

```
__cxx_global_var_init.exit:                                ; preds = %0  
%8 = bitcast i8* %1 to i32*  
store i8* %1, i8** bitcast (%"class.std::vector"* @bar to i8**)   
%9 = getelementptr inbounds i8* %1, i64 64  
store i8* %9, i8** bitcast (i32** getelementptr inbounds (%"class.std:  
store i32 0, i32* %8  
%10 = getelementptr inbounds i8* %1  
%11 = bitcast i8* %10 to i32*  
store i32 1, i32* %11  
%12 = getelementptr inbounds i8* %1  
%13 = bitcast i8* %12 to i32*  
store i32 2, i32* %13  
%14 = getelementptr inbounds i8* %1  
%15 = bitcast i8* %14 to i32*
```

Iterating over a Vector

```
double foo(std::vector<double> const& v) {  
    double s = 0.;  
    for(auto dat : v)  
        s += dat;  
    return s;  
}  
  
double bar(double const* v, std::size_t n) {  
    double s = 0.;  
    for(std::size_t i = 0; i < n; ++i)  
        s += v[i];  
    return s;  
}
```

Iterating over a vector

Auto version

```
.lr.ph:                                     ; preds = %.lr.ph.preheader
  %s.01 = phi double [ %8, %.lr.ph ], [ 0.000000e+00, %.lr.ph.preheader ]
  %6 = phi double* [ %9, %.lr.ph ], [ %2, %.lr.ph.preheader ]
  %7 = load double* %6
  %8 = fadd double %s.01, %7
  %9 = getelementptr inbounds double* %6, i64 1
  %10 = icmp eq double* %9, %4
  br i1 %10, label %._crit_edge.loopexit, label %.lr.ph
```

Iterating over a vector

Index version

```
.lr.ph:  
  %i.02 = phi i64 [ %5, %.lr.ph ], [ 0, %.lr.ph.preheader ]  
  %s.01 = phi double [ %4, %.lr.ph ], [ 0.000000e+00, %.lr.ph.preheader ]  
  %2 = getelementptr inbounds double* %v, i64 %i.02  
  %3 = load double* %2  
  %4 = fadd double %s.01, %3  
  %5 = add nuw i64 %i.02, 1  
  %exitcond = icmp eq i64 %5, %n  
  br i1 %exitcond, label %._crit_edge.loopexit, label %.lr.ph
```

Indexing a vector

```
double foo(std::vector<double> const& data, std::size_t i) {  
    return data[i];  
}  
double bar(double const* data, std::size_t i) {  
    return data[i];  
}
```

Indexing a vector

Needs an extra adress computation

```
define double @_Z3fooRKSt6vectorIdSaIdEEm(%"class.std::vector"* nocapture  
    %1 = getelementptr inbounds %"class.std::vector"* %data, i64 0, i32 0,  
    %2 = load double** %1  
    %3 = getelementptr inbounds double* %2, i64 %i  
    %4 = load double* %3  
    ret double %4  
}  
  
define double @_Z3barPKdm(double* nocapture readonly %data, i64 %i) {  
    %1 = getelementptr inbounds double* %data, i64 %i  
    %2 = load double* %1  
    ret double %2  
}
```

New / Delete

```
double foo(double x, double y) {  
    auto * z = new double( x + y);  
    auto tmp = *z;  
    delete z;  
    return tmp;  
}
```


New / Delete

No more heap allocation! (what about exceptions?)

```
define double @_Z3foodd(double %x, double %y) {  
    %1 = fadd double %x, %y  
    ret double %1  
}
```

Except / NoExcept

```
int bar() noexcept;  
int foobar();  
int foo() {  
    try {  
        return bar() + foobar();  
    }  
    catch(...) {  
        throw;  
    }  
}
```

Except / NoExcept

No global registration

```
declare i32 @_Z3barv() #1
```

```
declare i32 @_Z6foobarv() #2
```

```
attributes #1 = { nounwind "less-precise-fpmad"="false" "no-frame-pointer-elim"="false" }
```

```
attributes #2 = { "less-precise-fpmad"="false" "no-frame-pointer-elim"="false" }
```

Except / NoExcept

Different calling mechanism!

```
define i32 @_Z3foov() #0 {
    %1 = tail call i32 @_Z3barv()
    %2 = invoke i32 @_Z6foobarv()
        to label %3 unwind label %5

; :3                                     ; preds = %0
    %4 = add nsw i32 %2, %1
    ret i32 %4

; :5                                     ; preds = %0
    %6 = landingpad { i8*, i32 } personality i8* bitcast (i32 (...)* @__gxx_personality_v0) to:
        catch i8* null
    %7 = extractvalue { i8*, i32 } %6, 0
    %8 = tail call i8* @__cxx_b__begin_catch(i8* %7)
    invoke void @__cxx_b__rethrow()
        to label %15 unwind label %9

; :9                                     ; preds = %5
    %10 = landingpad { i8*, i32 } personality i8* bitcast (i32 (...)* @__gxx_personality_v0) to:
        cleanup
    invoke void @__cxx_b__end_catch(i8* %10)
        to label %11 unwind label %12
```

[illegible]

Outer Product

```
std::vector<double> outer_product(std::vector<double> const& self, std::v
    std::vector<double> outer(self.size() * other.size());
    for(std::size_t i = 0; i < self.size(); ++i)
        for(std::size_t j = 0; j < other.size(); ++j)
            outer[i * other.size() + j] = self[i] * other[j];
    return outer;
}
```

Outer Product

Vectorized (compiled with -O3 -march=native)

```
%wide.load20 = load <4 x double>* %76  
%77 = fmul <4 x double> %54, %wide.load  
%78 = fmul <4 x double> %59, %wide.load18  
%79 = fmul <4 x double> %63, %wide.load19  
%80 = fmul <4 x double> %68, %wide.load20
```

Bring Home Message

♥ COMPILERS ♥

They do their job, we do ours

Learn to communicate with them

`info gcc & clang --help`

☺ **trust no one & be curious** ☺

Verify assembly, verify code

THE END

HIGH-LEVEL THANKS

Joël Falcou and Pierrick Brunet

LOW-LEVEL THANKS

Adrien Guinet and Kévin Szudlapski

ALL-AROUND THANKS

CppCon & Quarkslab for allowing me to be there!