



ITERATOR HAIKU

HOW THE RANGES TS TURNED 5 ITERATOR CATEGORIES INTO 7,
AND BACK INTO 5 AGAIN

CASEY@CARTER.NET

OVERVIEW

- Ranges and iterators in Standard C++
- Iterator categories in the Palo Alto TR
- Ranges and iterators in the Ranges TS
 - Concept design flaws
 - Investigation leading to inconsistency
 - Re-design

RANGES IN STANDARD C++

- Why do we need a TS if the Standard already has ranges?!?
 - I have a family to feed
- Standard has **two** kinds of ranges

RANGE-SEQUENCE: ELEMENTS "BETWEEN" TWO ITERATORS

- Often denoted $[i, k)$
- Implicitly requires reachability
- `==` is an end-of-sequence test

RANGE-OBJECT: THINGS THAT WORK IN RANGE-BASED-FOR

- “Range” is right there in the name
- `begin(thing)` and `end(thing)` return iterators
 - `[begin(thing), end(thing))`

THE FIVE C++14 ITERATOR CATEGORIES

- Input
 - `*`, `++`, `==`, `!=`
 - Single-pass
 - Readable
 - Writable?

```
auto foo = ifstream{"integers.txt"};  
using I = istream_iterator<int>;  
assert(is_sorted(I{foo}, I{}));
```

THE FIVE C++14 ITERATOR CATEGORIES

- Forward
 - Input
 - Multi-pass

```
auto foo = forward_list<int>{0,1,2,3};  
assert(is_sorted(begin(foo), end(foo)));
```

THE FIVE C++14 ITERATOR CATEGORIES

- Bidirectional
 - Forward
 - --

```
auto foo = list<int>{0,1,2,3};  
reverse(begin(foo), end(foo));
```


THE FIVE C++14 ITERATOR CATEGORIES

- RandomAccess
 - Bidirectional
 - -, +, -, +=, -=, []

```
auto foo = vector<int>{0,1,2,3};  
sort(begin(foo), end(foo));
```

THE FIVE C++14 ITERATOR CATEGORIES

- Output
 - `*`, `++`
 - Single-pass
 - Writable
 - Readable?

```
auto foo = vector<int>{0,1,2,3};  
using O = ostream_iterator<int>;  
copy(begin(foo), end(foo), O{cout});
```

DOMAIN OF A FUNCTION

- The set of values over which the function (operation, expression) is defined

```
int f(int i) {  
    return 42 / i;  
}
```

FORWARD: DOMAIN OF ==

- Iterators over the same underlying sequence (N4606 [forward.iterators]/2)

OUTPUT: DOMAIN OF ==

- There is no spoon

INPUT: DOMAIN OF $==$

... the term *the domain of* $==$ is used in the ordinary mathematical sense to denote the set of values over which $==$ is (required to be) defined. This set can change over time. Each algorithm places additional requirements on the domain of $==$ for the iterator values it uses. These requirements can be inferred from the uses that algorithm makes of $==$ and $!=$. [*Example*: the call `find(a, b, x)` is defined only if the value of `a` has the property p defined as follows: `b` has property p and a value `i` has property p if $(*i == x)$ or if $(*i != x$ and $++i$ has property p). —end example]

PALO ALTO TR

- N3351 “A Concept Design for the STL” (<http://wg21.link/n3351>)
- Redefine iterators and algorithms in terms of Concepts
 - Replaces tables in the spec with syntax in the library
- Awesome. But what’s a concept?

CONCEPTS

- Syntax (Concepts TS):
 - Predicates over a set of parameters
 - Requirements on types (e.g., “auto foo = bar;”)
 - Associated types that must exist
 - Expressions that must be valid
- Semantics (Ranges TS):
 - Requirements on values (e.g., “foo must equal bar”)

PALO ALTO TR

- Some algorithms take unpaired input iterators
 - “3-legged” double-range algorithms
 - `equal(first1, last1, first2)`
 - iterator + count algorithms
 - `copy_n(first, count, out)`
- Unpaired input iterators don't need `==` and `!=`
 - New category `WeakInputIterator`

RANGES TS

- Range-object algorithms
 - Turn a range-object into a range-sequence with begin and end
 - Invoke the range-sequence algorithm
- Output ranges?
- Sentinels?
 - Not the X-Men hunting killer robots
 - A new way to denote range-sequences

RANGES TS: OUTPUT

- Rename OutputIterator to WeakOutputIterator
- Add a new OutputIterator with == and !=
 - WeakOutput : Output == WeakInput : Input

5 HAVE BECOME 7

- Output, Input, Forward, Bidirectional, RandomAccess, **WeakOutput**, **WeakInput**
- Ranges - now with 40% more categories!

ITERATOR VS. WEAKITERATOR??

- Factor commonality out of WeakInput and WeakOutput into WeakIterator
- Factor commonality out of Input and Output into Iterator
- "iterators" colloquially means....models of "WeakIterator"?!?

INTRODUCE SENTINELS INTO THE MODEL

- A distinct *value* that denotes the end of a range
 - `istream_iterator<T>()`
- Why not use a distinct *type*?

SENTINEL CONCEPT

- Relates a type `I` that satisfies `Iterator` and a type `S` that satisfies `Regular`
- Requires `I` and `S` to satisfy `EqualityComparable`

CROSS-TYPE EQUALITYCOMPARABLE

- Cross-type relations require relation to hold for
 - Both individual types
 - Common type
 - What's a common type?

COMMON TYPE

- C is a common type of T and U if:
 - $C(t)$ is valid
 - $C(t1) == C(t2)$ iff $t1 == t2$
 - (symmetric requirement for u)
- e.g.:
 - `int` is a common type of `int` and `short` (so is `long`, and `long long`)
 - `int` is a common type of `int` and `int`
 - `std::string` is NOT a common type of `char*` and `char const*` (equality of pointers is different from equality of pointers-converted-to-strings)

CROSS-TYPE EQUALITYCOMPARABLE

- Cross-type relations require relation to hold for
 - Both individual types
 - Common type
- Cross-type EC also requires that:
 - $C(t1) == C(u) \ \&\& \ t1 == t2 \text{ implies } C(t2) == C(u)$
 - (symmetric requirement)

In other words, the equality is transitive across types.

SENTINEL CONCEPT

- Relates a type `I` that satisfies `Iterator` and a type `S` that satisfies `Regular`
- Requires `I` and `S` to satisfy `EqualityComparable`
 - `I` and `S` must be individually `EC`
 - No “weak” ranges
 - despite that algorithms don’t use `==` on sentinels or single-pass iterators

ALGORITHMS DON'T USE == FOR SINGLE-PASS ITERATORS

```
template<InputIterator I, Sentinel<I> S, /*...*/>
bool any_of(I first, S last, P predicate) {
    for (; first != last; ++first) {
        if (invoke(predicate, *first)) {
            return true;
        }
    }
    return false;
}
```

ALGORITHMS DON'T USE == FOR SINGLE-PASS ITERATORS

- Concepts require == / != anyway
- Either useless or have undefined behavior
 - Not just a footgun, but a *required* footgun!

SENTINEL CONCEPT

- Relates a type `I` that satisfies `Iterator` and a type `S` that satisfies `Regular`
- Requires `I` and `S` to satisfy `EqualityComparable`
 - `I` and `S` must be individually `EC`
 - No “weak” ranges
 - despite that algorithms don’t use `==` on sentinels or single-pass iterators

ITERATORS AND SENTINELS ARE BOTH “POSITIONS”

- General recommendation to treat all sentinels as equal
 - Sentinel represents the “at the end” position

AXIOMATIZATION OF ITERATORS WITH SENTINELS

- Throw out everything to do with ranges-as-iterator pairs
- Start over with only $[\text{iterator}, \text{sentinel})$ ranges
- Derive range properties
- Define domains of iterator and sentinel operations)
- Sadly not enough time to cover here
 - Cool math like “if $[i, s)$ denotes a range then either $i == s$ or $[++i, s)$ denotes a range”
 - Everyone likes derivations
- Woops, problem with stateful sentinels

STATEFUL SENTINELS AND “ALWAYS EQUAL”

```
struct S {  
    int i;  
  
    friend bool operator==(int const* p, S const& s) {  
        return *p >= s.i;  
    }  
    friend bool operator==(S const&, int const*);  
    friend bool operator==(S const&, S const& ) {  
        return true;  
    }  
    // define operator!= overloads appropriately  
};
```

STATEFUL SENTINELS AND “ALWAYS EQUAL”

```
int a[] = {2, 1, 3};

assert(a+1 != S{2}); // !(1 >= 2)
assert(a+2 == S{2}); // (3 >= 2), so [a+1,S{2}) is {1}

assert(a+1 != S{3}); // !(1 >= 3)
assert(a+2 == S{3}); // (3 >= 3), so [a+1,S{3}) is {1}
assert(S{2} == S{3}); // Cross-type EC

assert(a+0 == S{2}); // (2 >= 2), so [a+0,S{2}) is {}
assert(a+0 == a+2); // Cross-type EC
```

S SURE LOOKS LIKE A SENTINEL TO ME

- “Sentinels are positions” vs “Sentinels are predicates”
- How did this problem not show up in the two implementations?

START OVER WITH ALGORITHM REQUIREMENTS

- Algorithms care about:
 - $i == s, s == i, i != s, s != i$ have the same domain
 - Symmetry: $(i == s) == (s == i)$ and $(i != s) == (s != i)$
 - Complement: $i != s == !(i == s)$
 - $i == s$ is well-defined when $[i, s)$ denotes a range
- First three look like general equality comparison requirements
 - Define a weaker cross-type equality for iterators and sentinels

WEAKLYEQUALITYCOMPARABLE

- Non-transitive cross-type equality requirement for iterators/sentinels
- Prior EqualityComparable concepts refine WeaklyEqualityComparable
- Yuck: invoking a predicate with `==` syntax
 - Semantic abomination
 - Transitivity is absent, but preserves other EC semantics
 - Backwards compatible

WEAK RANGES ARE POSSIBLE

- Without cross-type `EqualityComparable`, no need for individually `EqualityComparable` types
- Sentinel becomes:
 - Relates a type `I` that satisfies `WeakIterator` and a type `S` that satisfies `Semiregular`
 - Requires `I` and `S` to satisfy `WeaklyEqualityComparable`

STATEFUL SENTINELS WORK

```
struct S {  
    int i;  
  
    friend bool operator==(int const* p, S const& s) {  
        return *p >= s.i;  
    }  
    friend bool operator==(S const&, int const*);  
    friend bool operator==(S const& x, S const& y) {  
        return x.i == y.i;  
    }  
    // define operator!= overloads appropriately  
};
```

STATEFUL SENTINELS WORK

```
int a[] = {2, 1, 3};

assert(a+1 != S{2}); // !(1 >= 2)
assert(a+2 == S{2}); // (3 >= 2), so [a+1,S{2}) is {1}

assert(a+1 != S{3}); // !(1 >= 3)
assert(a+2 == S{3}); // (3 >= 3), so [a+1,S{3}) is {1}
assert(S{2} == S{3}); // Cross-type EC
assert(S{2} != S{3});

assert(a+0 == S{2}); // (2 >= 2), so [a+0,S{2}) is {}
assert(a+0 == a+2); // Cross-type EC
```


BUT THEN "STRONG" CONCEPTS AREN'T NEEDED

- A type satisfies `(Input | Output |)Iterator` if it
 - satisfies `Weak(Input | Output |)Iterator`
 - is a sentinel for itself (i.e., `T` satisfies `Sentinel<T, T>()`)
- Algorithms don't need "strong" anyway
- If we don't have "strong" variants anymore, why use the "Weak" prefix?

7 BECOMES 5

- Augment (InputIterator | OutputIterator | Iterator) requirements with Sentinel requirements
- Strip “Weak” prefixes
- Replaces the “weak vs strong” distinction with “sentinel vs non-sentinel”

MOST IMPORTANTLY

- Writers of single-pass iterators and sentinels need not write:
 - `bool operator == (..) { return true; }`
- We've turned undefined behavior into a compile error!
 - Well, sort of.
- We've reduced the number of concepts
- “Iterators” actually means iterators

TAKEAWAY: CONCEPT LIBRARY DESIGN

- Concepts that don't exactly fit the usage requirements smell
 - Trade-off *Minimality* with *Purity*
- Deriving requirements from first principles can give clarity
 - I did not see the inherent contradiction in:
 - Sentinels are always equal
 - Iterators are sometimes Sentinels
 - Iterators are never always equal

HOW CAN I LEARN MORE?

- Ranges TS working paper: <http://wg21.link/n4569>
- Early design blogs at <http://ericniebler.com/>
- Implementation with Concepts: <https://github.com/CaseyCarter/cmcstl2> (GCC 6+)
- Implementation with C++11 “concepts”: <https://github.com/EricNiebler/range-v3> (GCC 4.8+, Clang 3.3-ish, Clang/C2)
- Implementation with extensive MSVC workarounds:
<https://github.com/Microsoft/Range-V3-VS2015> VS2015 Update 3

LIBRARIES ACQUISITION

- 80% of C++ projects use 2 or more 3rd party libs
 - A majority of them use open source libraries
- Acquiring and rebuilding libs on Windows can be simple
 - NuGet wasn't designed for C++ (e.g. no local rebuilding)
- Open source tool based on a port tree approach: "VCPKG"
 - Usage: VCPKG install Boost
 - Installs the .h, .lib and binaries in a "lib folder" ready to use

<https://github.com/Microsoft/VCPKG>

boost
cpprestsdk
curl
expat
freetype
glew
glfw3
libjpeg-turbo
libpng
libuv
libwebsockets
mpg123
openal-soft
opencv
openssl
range-v3
SDL2
sqlite3
tiff
tinyxml2
zlib