

# Out of memory? Business as usual.

Examining applications that do not terminate on `std::bad_alloc`

```
Sep 19 00:33:10 server1 kernel: Out of Memory: Killed process 8631 (xterm).
Sep 19 00:33:34 server1 kernel: Out of Memory: Killed process 9154 (xterm).
Sep 19 00:34:05 server1 kernel: Out of Memory: Killed process 6840 (xterm).
Sep 19 00:34:42 server1 kernel: Out of Memory: Killed process 9066 (xterm).
Sep 19 00:35:15 server1 kernel: Out of Memory: Killed process 9269 (xterm).
Sep 19 00:35:43 server1 kernel: Out of Memory: Killed process 9351 (xterm).
Sep 19 00:36:05 server1 kernel: Out of Memory: Killed process 21188 (sshd).
```



Aw,

Something went wrong while displaying this webpage.

[Learn more](#)

Reload

Software exception Detected

exception, it will now attempt to save all  
unsaved documents and exit.

Exception: std::bad\_alloc

OK

LibreOffice 5.

std::bad\_alloc

OK

Out of memory attempting to open smalltext.txt

Print Preview Error



There is not enough free memory to print.

OK

## Page crashed

Unfortunately, something caused this page to quit. It might have  
been an extensions conflict or some other reason.

**Try reloading the page, or navigate to another page to  
continue.**

Reload this page

Oops!

Something went wrong while displaying this page. Please  
reload or visit a different page to continue.

Reload

Writer



Errors occur when WPS Writer is opening this file.  
Please try one of the following.

- Make sure you are permitted to access this file or drive.
- Make sure there is enough memory and disk space.

OK

Sergey Zubkov

# What is this about?

1. What is “out of memory” and what is “bad allocation”
2. What do those who catch `std::bad_alloc` do with it
3. How common (or rare) are these applications

# All resources are not created equal

- “Resource” == “limited availability” (Wikipedia, cppreference)
- Assumed resources
  - CPU time/cores/caches, Network bandwidth, RNG entropy, Electricity, **Stack memory**
- Checked resources
  - Disk space, other hardware, file/socket descriptors, threads, locks, other software, **Heap memory**

Resource acquisition can fail: constructors can throw. What do they throw in stdlib?

- `std::system_error`:
  - `thread`, `unique_lock`, `shared_lock`, `maybe_shared_ptr`
  - `basic_socket` (Network TS), `display_surface` (Graphics TS)
- `std::bad_alloc`:
  - `any`, `shared_ptr`, `function`, `boyer_moore_(horspool_)searcher`, `basic_string`, `(forward_)list`, `vector`, `deque`, `(unordered_)(multi)(map|set)`, `stack`, `(priority_)queue`, `valarray`, `basic_(i|o)stringbuf`, `stringstream`, `(un)synchronized_pool_resource`, `path`, `directory_entry`, `(runtime_|logic_|etc)error`, `basic_regex`, `match_results`, `promise`, `packaged_task`, ...

# Common wisdom

- Effective C++ (2<sup>nd</sup> ed. Item 7, 2003, no longer in 3<sup>rd</sup> ed.)  
“Regardless of whether you use "normal" (i.e., exception-throwing) new or "nothrow" new, it's important that you be prepared to handle memory allocation failures.”
- Sutter’s Mill “To new, perchance to throw part 2”, 2001  
“except for special cases, even when you detect new failure there's not always much you can do if there really is no memory left.”
- DIP-33 rationale by Walter Bright, 2013  
“I've almost never seen a program that could successfully recover from out of memory errors, even ones that purport to.”
- CppCoreGuidelines rule F.6, 2016  
“after memory runs out it is hard to do anything clever [...] the majority of programs and execution environments cannot meaningfully handle a failure to allocate”

# What does “Out of memory” mean?

- “Memory” means page-based virtual memory.  
Process sees a homogeneous address space, but each page may be private or shared, read-only or (COW-)writeable, clean or dirty, resident or paged-out
- Unused memory is wasted memory  
Operating systems swap or reclaim inactive pages for buffers/caches.
- *Commit charge*: all writeable pages that are not file-backed (stack, data, heap, private mmap, shared library .GOT's)
- If commit charge exceeds free RAM + free swap, it's OOM... or is it?

# The fork/exec problem

- fork() duplicates the entire process memory, exec() throws it away

```
$ ./test 10
```

```
allocating 10737418240 bytes (10GB)...Allocated
```

```
parent: fork ok
```

```
child: fork ok
```

```
$ ./test 20
```

```
allocating 21474836480 bytes (20GB)...Allocated
```

```
fork: Not enough space
```

- Not an issue on Windows
- Alternatives (vfork and spawn) exist, but have their own issues

# Overcommit

## Operating systems with strict commit accounting:

- Windows, Solaris, HP/UX, and more;

Note: private COW mappings are fully accounted; fork() can OOM

## Operating systems with overcommit and OOM killer:

- AIX: per-process opt-out by installing SIGDANGER handler
- FreeBSD: can be turned off systemwide, process can opt-out with protect(1)
- Linux: three systemwide settings: always, never, and heuristic (+oom\_adj, +cgroups).

## Some Linux users swear by it, some swear it off

Note: in “always overcommit” mode, rlimits and address space limits will still throw

Note: in “never overcommit”, kernel reserves RAM to fork a shell/top/kill

Note: neither mode is the default: the default is heuristic



# Common ~~wisdom~~ myths

Sutter's Mill "To new, perchance to throw part 2" (and many others)

"checking for new failure isn't as important as one might think ... On some operating systems, including specifically Linux, memory allocation always succeeds."

LevelDB issue #335 "Question about exception safety"

"on Linux, you will only get a `std::bad_alloc` thrown if the virtual address space has been exhausted"

Note: the Linux default is not always-overcommit

# Bad allocation is not always OOM

This is a bad allocation:

```
std::vector<int>(-1); // always throws (including Linux with overcommit_always)
```

That's clearly programmer's fault. The first unit test will catch it. Why care?

```
std::vector<int>(content_length); // do you trust Content-Length:?
```

How hard can it be to check your inputs?

- CVE-2016-2109 OpenSSL OOM DoS due to short invalid encoding
- CVE-2016-2463 Android OOM DoS via crafted media file
- CVE-2016-6170 ISC BIND OOM DoS due to large UPDATE message
- CVE-2015-7540 samba AD-DC OOM DoS via crafted packets
- CVE-2015-1819 libxml OOM DoS via crafted XML file
- CVE-2014-3506 OpenSSL OOM DoS via crafted DTLS handshake
- CVE-2013-7447 cairo OOM DoS via crafted image file

# Aren't all those CVEs against C libs?

`malloc` returns NULL, but how do you inform the caller, possibly many stack frames up?

- Some meticulously return error codes
- Some `longjmp`
- Some gave up:

```
mem = malloc (n_bytes);  
TRACE (GLIB_MEM_ALLOC((void*) mem, (unsigned int) n_bytes, 0, 0));  
if (mem)  
    return mem;  
  
g_error ("%s: failed to allocate %s"G_GSIZE_FORMAT" bytes",  
        G_STRLOC, n_bytes);
```

<https://github.com/GNOME/glib/blob/master/glib/gmem.c#L87-L106>

# What would C++ do?

```
if (reinterpret_cast<const uint8_t*>(&_isize) == buffer) {  
    // Read length of data to follow  
    try {  
        _isize = ntohl(_isize);  
        if (0 == _isize || _isize > MAX_XRL_INPUT_SIZE)  
            throw bad_alloc();  
        _input_buffer.resize(_isize);  
    } catch (bad_alloc) {  
        XLOG_ERROR("Bad input buffer size (%d bytes) from wire, "  
                    "dropping connection", XORP_INT_CAST(_isize));  
        error_event();  
        return;  
    }  
}
```

[https://github.com/greearb/xorp.ct/blob/master/xorp/libxipc/finder\\_tcp.cc#L161-L173](https://github.com/greearb/xorp.ct/blob/master/xorp/libxipc/finder_tcp.cc#L161-L173)

# “One Bug To Rule Them All”

## (GSEC-TZO-44-2009 aka CVE-2009-1692)

IE5, IE6, IE7, IE8, Netscape, Firefox, Safari, Opera, Konqueror, Seamonkey, Wii, PS3, iPhone, iPod, Nokia, Siemens.... and more

Konqueror (Ubuntu): allocates 2GB of memory then either crashes the browser or (most often) the OS reboots.

Chrome :allocates 2GB of memory then crashes tab

Firefox : allocates 2GB of memory then the browser crashes

IE5,6,7,8 : allocates 2GB of memory then the browser crashes

Opera : will not crash but other applications will become unstable

Nintendo WII (Opera) : Console hangs, needs hard reset

Sony PS3 - Console hangs, needs hard reset

iPhone - iPhone hangs and needs hard reset

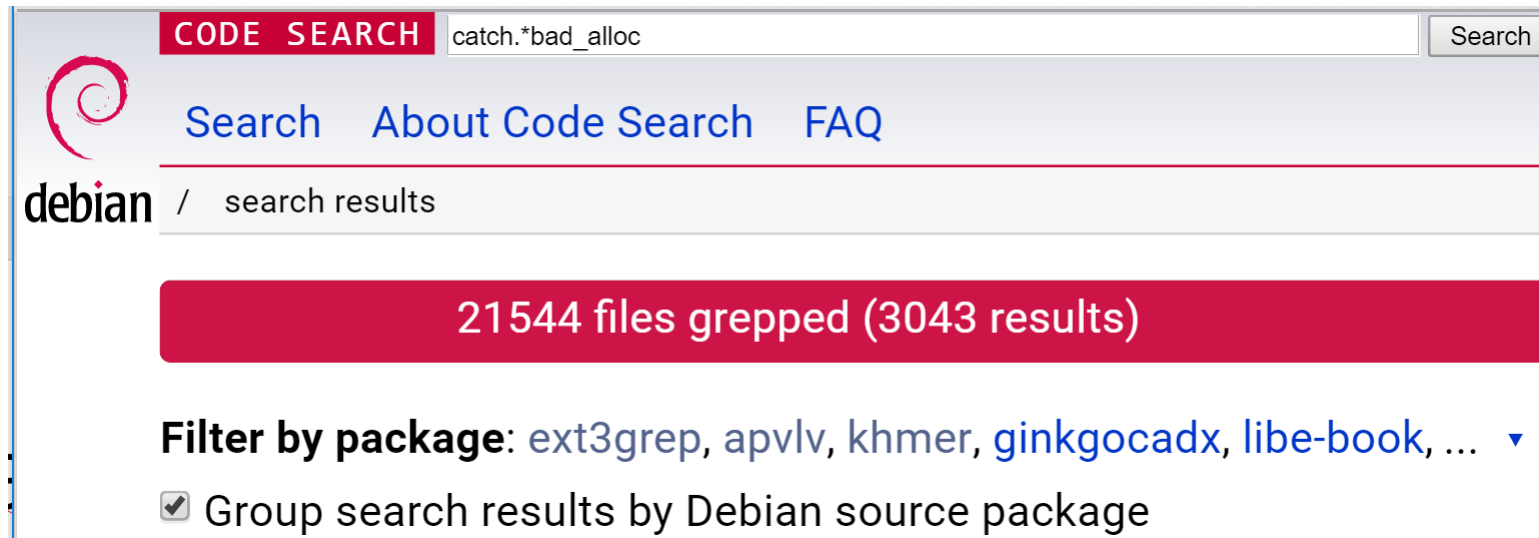
What was this awful bug?

```
e = document.createElement("select");  
e.length=2147483647;
```

Just one of infinite ways by which JavaScript can use memory

# Does anyone really handle `std::bad_alloc`?

Debian Code Search: `codesearch.debian.net/search?q=catch.*bad_alloc`



The screenshot shows the Debian Code Search interface. At the top, there is a search bar with the text "CODE SEARCH" and a search button. The search query "catch.\*bad\_alloc" is entered in the search bar. Below the search bar, there are links for "Search", "About Code Search", and "FAQ". The Debian logo is visible on the left. The search results section shows "21544 files grepped (3043 results)". Below this, there is a filter by package section with a dropdown menu showing "ext3grep, apv, khmer, ginkgocadx, libe-book, ...". A checkbox is checked for "Group search results by Debian source package".

Not `std::exception` or `catch(...)`, only explicit handling for `std::bad_alloc`:

3043 results in 341 packages.

Disclaimer: I can't promise I fully understood each of those 3000+ pieces of opensource code

# 178 (46%) Somebody else's problem

92 (23%) Convert to error code (either returned or stored as a flag)

gdal qscintilla libsdl2 rtaudio armadillo libstdc++ breakpad ACE vtk coinutils ...

50 (13%) Convert to custom exception

dlib ipopt poco povray gdal libreoffice capnproto galera-3 libosl gnu datalanguage ...

20 (5%) Convert to OOM error in a different language (PyErr\_NoMemory)

boost cython pyclingo matplotlib openjdk libreoffice openvrml healpy pytaglib ...

16 (4%) Rethrow as-is

mp3diags 3dldf timbl deqp cxxtools orthanc-postgresql duma synergy mrpt madness ...

# Somebody else's problem

## 92 (23%) Convert to error code

Many libraries return error codes

```
try {  
    int position = Length();  
    InsertString(position, data, length);...  
} catch (std::bad_alloc &) {  
    return SC_STATUS_BADALLOC;
```

<https://github.com/mirror/scintilla/blob/master/src/Document.cxx#L1056-L1060>

the caller can handle this generically:

```
if (_pscratchTilla->execute(SCI_GETSTATUS) != SC_STATUS_OK)  
    throw;
```

<https://github.com/notepad-plus-plus/notepad-plus-plus/blob/master/PowerEditor/src/ScitillaComponent/Buffer.cpp#L1423-L1424>

## Some update global or member variables instead

```
__try  
{ __words = new (std::nothrow) _Words[__newsize]; }  
__catch(const std::bad_alloc&)  
{ __words = nullptr; }  
if (!__words)  
{  
    _M_streambuf_state |= badbit;
```

<https://github.com/gcc-mirror/gcc/blob/master/libstdc++-v3/src/c++11/ios.cc#L127-L133>



# Somebody else's problem

## 50 (13%) Convert to custom exception

```
try
{
    if (aPosition >= _bindVector.size())
    {
        _bindVector.resize(aPosition + 1);
    }

    InputParameter inputParameter(aFieldType, aBufferPtr, aLength);

    _bindVector[aPosition] = inputParameter;
}
catch (std::bad_alloc&)
{
    PostgreSQLException("Memory allocation error while binding");
}
```

# Somebody else's problem

## 16 (4%) Rethrow as-is

### Clean up non-RAII resources

```
        m_allocations.push_back(ptr);
    }
    catch (std::bad_alloc& )
    {
        delete [] ptr;
        throw;
    }
```

[https://github.com/zenoalbisser/chromium/blob/master/third\\_party/deqp/src/framework/referencerenderer/rrVertexPacket.cpp#L64-L72](https://github.com/zenoalbisser/chromium/blob/master/third_party/deqp/src/framework/referencerenderer/rrVertexPacket.cpp#L64-L72)

### Handle other exceptions differently

```
    catch( const std::bad_alloc& e )
    {
        throw e;
    }
    catch(...)
    {
        _allocator.deallocate(_instance, 1);
        _instance = 0;
        throw;
    }
```

<https://github.com/deniskin82/cxxtools/blob/master/include/cxxtools/singleton.h#L85-L94>

# 82 (21%) Cleanup and terminate

- 45 (12%) Not from main  
rethinkdb ipopt fluxbox lzip shogun thrift polyml krita pingus bowtie ...
- 32 (8%) From main  
tripwire smartmontools tango timbl taskd ppl ossim mame dwarfutils ...
- 3 (1%) Configurable handler defaulting to abort  
igraph r-cran-igraph gdal
- 1 (0%) “Parachute” (as featured in Code Complete)  
scantailor

# 82 (21%) Cleanup and terminate

“Cleanup” often means removing temp or lock files

```
} catch (const std::bad_alloc &) {  
    unlink_ofile(oname);  
    printErr(iname, "out of memory");  
    e_exit(EXIT_ERROR);  
}
```

<https://github.com/ferseiti/upx-ucf/blob/master/src/work.cpp#L310-L313>

Termination from deep in a library may have major impact:

```
void* fastMalloc(size_t n)  
{  
    void* result = malloc(n);  
    if (!result)  
        CRASH();  
}
```

<https://github.com/WebKit/webkit/blob/master/Source/WTF/wtf/FastMalloc.cpp#L131-L135>

# 35 (9%) Moving forward

- 10 (3%) Try allocating less (smaller audio buffer, on-the-fly calcs)  
audacity eigen3 frobby tuvok lammips libreoffice mira openttd spring vxl
- 10 (3%) Swallow bad\_alloc (in a destructor, adding to cache)  
frobby jade libreoffice love poco ncbi-blast+ opencv openvrml ossim otb
- 7 (2%) Alternative algorithm (in-place instead of out-of-place a-la STL)  
vtk aseprite bowtie bowtie2 krita mlpack octave
- 6 (2%) Free up some memory (drop caches, cannibalize freelists)  
libstdc++ mrpt sonic-visualizer scylladb diagnostics libosl
- 2 (1%) Just try again (???)  
gnuradio infinidb

# 35 (9%) Moving forward

## 7 (2%) In-place algorithm

```
bool inline inplace_transpose(arma::Mat<eT>& X)
{
    try
    {
        X = arma::trans(X);
        return false;
    }
    catch (std::bad_alloc&)
    {
        #if (ARMA_VERSION_MAJOR >= 4) || \
            ((ARMA_VERSION_MAJOR == 3) && (ARMA_VERSION_MINOR >= 930))
            arma::inplace_trans(X, "lowmem");
        return true;
    }
}
```

[https://github.com/mlpack/mlpack/blob/master/src/mlpack/core/data/load\\_impl.hpp#L65-L77](https://github.com/mlpack/mlpack/blob/master/src/mlpack/core/data/load_impl.hpp#L65-L77)

# 68 (18%) Roll back and do something else

- 47 (12%) Interactive apps refusing user actions (“Could not open file”)  
notepad-plus-plus libreoffice lfhex textstudio inkscape spring povray ...
- 18 (5%) Servers dropping service requests
  - Network servers
    - apt-cacher-ng ntopng dc-qt dlib folly reciprocate xorp
  - Databases and other servers
    - clamav csound dindel glogg libclasp ring scylladb tarantool
  - Batch processors
    - clblas seqan ssdeep undertaker
- 3 (1%) Prepared fallback (Error texture, “NoSound” sound driver)  
Oad aiksaurus desmume

68 (18%) Roll back and do something else  
47 (12%) UI and other interactive apps

“File too big to load” is the most common reason for this

```
        docPtr = Poppler::Document::load(fileName);  
    }  
} catch (std::bad_alloc) {  
    error = PopplerErrorBadAlloc;  
    return QSharedPointer<Poppler::Document>();  
}
```

[TexStudio/blob/master/pdfrendermanager.cpp#L119-L123](#)

And plenty other reasons

```
catch(vtkstd::bad_alloc &)  
{  
    throw IRISException("Out of memory during mesh computation");  
}  
catch(IRISException & IRISexc)  
{  
    QMessageBox::warning(this, "Problem generating mesh", IRISexc.what());  
}
```

<https://github.com/pyushkevich/itksnap/blob/master/GUI/Model/Generic3DModel.cxx#L232-L234>



68 (18%) Roll back and do something else  
18 (5%) Servers dropping service requests

```
try {  
    pass_verdict = processPacket(&h->ts, time, ethernet, vlan_id, iph,  
                                ip6, h->caplen - ip_offset, h->len,  
                                h, packet, shaped, ndpiProtocol);  
} catch(std::bad_alloc& ba) {
```

<https://github.com/ntop/ntopng/blob/dev/src/NetworkInterface.cpp#L1393-L1397>

```
catch (std::bad_alloc &)  
{  
    std::cout << shapeStrings[i] << " threw bad_alloc exception, skipping this shape." << std::endl;  
    continue;  
}
```

<https://github.com/seqan/seqan/blob/master/apps/razers/paramChooser.h#L675-L679>

# 25 (6%) Unit tests

- Unit tests

gdcmlibstdc++ boost angle cmtk trilineos dune-common dune-istl eigen3 folly isc-kea  
libboost-geometry-utils-perl libcpp libloki libpqxx pugixml mia ossim ppl pugixml tbb vxl  
xapian-core z3 scylladb

```
ASSERT_THROWS_IN_TEST(  
    {  
        limit_foo_count_in_scope foo_limit(FooCount + planned_victim_size);  
        victim = std::move(fixture.source); // fragmented assignment  
    },  
    std::bad_alloc, "", test_name  
);
```

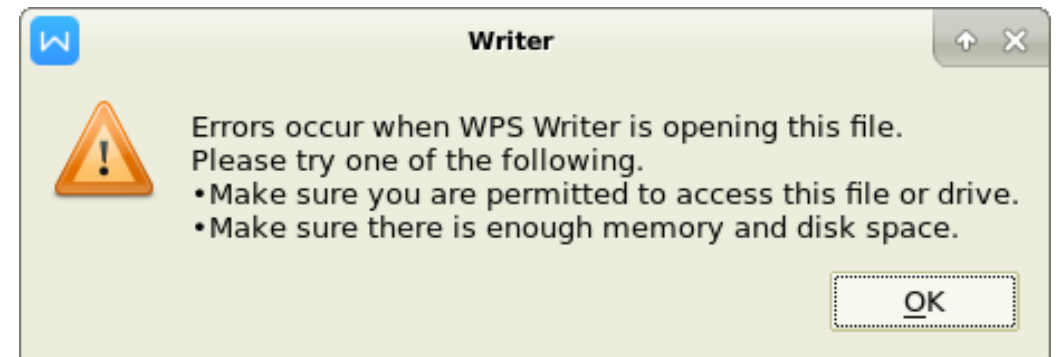
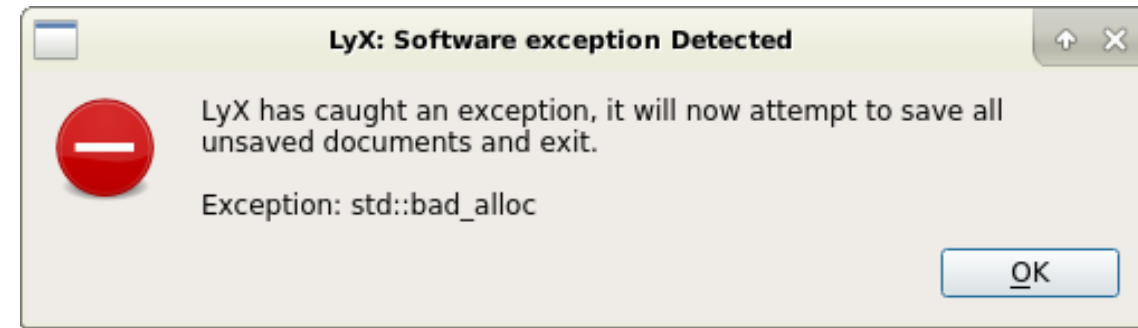
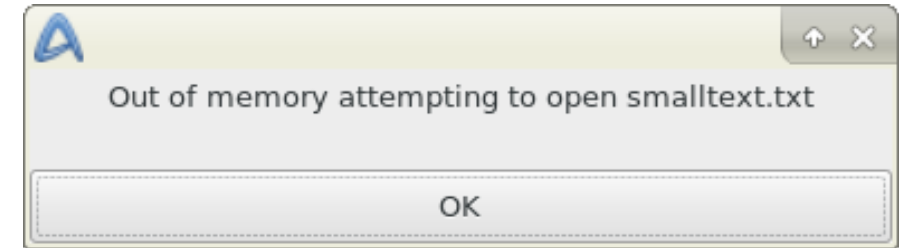
[https://github.com/wjakob/tbb/blob/master/src/test/test\\_concurrent\\_vector.cpp#L1407-L1413](https://github.com/wjakob/tbb/blob/master/src/test/test_concurrent_vector.cpp#L1407-L1413)

# Case studies or “crash everything”

We’ve seen what people do with `std::bad_alloc` when they catch it, but how common is that, really?

# C++ OOM Case study: Office applications

| App         | Large file       | Large paste            |
|-------------|------------------|------------------------|
| Abiword     | Survives         | Uncaught bad_alloc     |
| calligra    | Survives         | Survives, seems flaky  |
| lyx         | Save and quit    | Save and quit          |
| texmaker    | Uncaught         | Uncaught               |
| WPS office  | Survives         | Survived, then crashed |
| libreoffice | Save and survive | abort                  |
| openoffice  | Uncaught         | Segfault               |
| texstudio   | Qt handler abort | Survives               |
| texmacs     | Qt handler abort | Survives               |
| scribus     |                  | Uncaught               |
| qchartdiary |                  | Qt handler abort       |
| kmymoney    |                  | Qt handler abort       |
| kraft       |                  | Qt abort or size limit |



# C++ OOM Case study: Code Editors/IDEs

| App                    | Large file                        | Large paste   |
|------------------------|-----------------------------------|---|
| Sublime                | Quietly exits                     | Saves and quietly exits                             |
| Notepad++              | Survives because of scintilla     |   |
| Codeblocks             | Uncaught bad_alloc                | Glib crash or segv, scintilla doesn't save this one |
| Geany                  | Glib crash                        | Survives because of scintilla, menus etc all work.  |
| Leechcraft             | Quietly refuses to load, survives | Survives because of scintilla (but crashed once)    |
| U++ TheIDE             | Warning popup and exit            | Warning popup and exit                              |
| SciTE                  | Uncaught exception                | Uncaught exception                                  |
| Pikdev                 | Uncaught bad_alloc                | Survives, without Scintilla's help; no bad_alloc    |
| Sasm                   | Uncaught bad_alloc                | Qt handler crash                                    |
| Kdevelop, Monkeystudio | Qt handler crash                  | Qt handler crash                                    |
| Kate                   | Uncaught bad_alloc                | Uncaught bad_alloc                                  |
| Qt Creator             | segfault                          | Uncaught bad_alloc                                  |
| CodeLite               | Glib crash                        | Segfault  |
| Rstudio                | Imposes a file size limit         | Webkit crash  |

# C++ OOM Case study: Web browsers

| App  | Large page   | JavaScript OOM  |
|--|--|---|
| One tab per process:<br>Chromium, Epiphany,<br>Opera, Vivaldi,<br>Qupzilla | tab crash  | tab crash   |
| Mozilla family:<br>Firefox, Seamonkey,<br>Conkeror                         | Usually stops download,<br>but sometimes crashes in<br>glib or with “unhandleable<br>oom while tenuring” | Stops the script, but doesn’t<br>free (despite<br>SpiderMonkey’s two-level<br>OOM callbacks)                      |
| Webkit family:<br>QtWeb, dwb, Otter  | Webkit crash   | Stops the script, but doesn’t<br>free: further browsing<br>quickly gets webkit crash or<br>QThread::start failure |
| leechcraft, rekonq,<br>dillo, dooble                                       | crash  | crash   |



Aw, Snap!

Something went wrong while displaying this webpage.

[Learn more](#)

Reload

**Oops!**

Something went wrong while displaying this page. Please reload or visit a different page to continue.

Reload



**Page crashed**

Unfortunately, something caused this page to quit. It might have been an extensions conflict or some other reason.

**Try reloading the page, or navigate to another page to continue.**

Reload this page



# C++ OOM Case study: Databases

| App  |             |  |
|--|-------------|--|
| scylladb   | Survives... | except in future::then (open issue)                  |
| tarantool  | Survives... | except when lua OOMs                                 |
| rethinkdb, rocksdb, leveldb, kyoto, mongodb, etc |             | crash (some easier than others), but users complain: |

(Note: there are other databases that don't crash – this slide is incomplete)

<https://github.com/google/leveldb/issues/335>

“I found many palces in leveldb source code that bad\_alloc will break the Ref/Unref balance.”

<https://github.com/rethinkdb/rethinkdb/issues/599>

“A truly robust solution would deal with an out of memory condition with a response less drastic than killing the server, say by cancelling the offending query and freeing any memory associated with it”

# Handling allocation failure

- Users want it (even Linux users)
  - When data loss, expensive recalculation, or service disruption may occur
- Consistent RAll makes it possible
  - just don't leak another `bad_alloc` from a destructor
- Don't be perfect: heterogeneous memory handling strategies work
  - extreme case: preallocate everything
- Don't underestimate the importance of libraries
  - a library can break every user (glib, webkit)
  - or help every user (scintilla)