

The strange details of `std::string` at Facebook

Nicholas Ormrod

Software Engineer

CppCon 2016

Questions I have answers for

- What different string implementations exist?
- How are strings optimized?
- What goes wrong when hunting for improvements?


Questions I want answers for

- What is the most efficient string implementation?


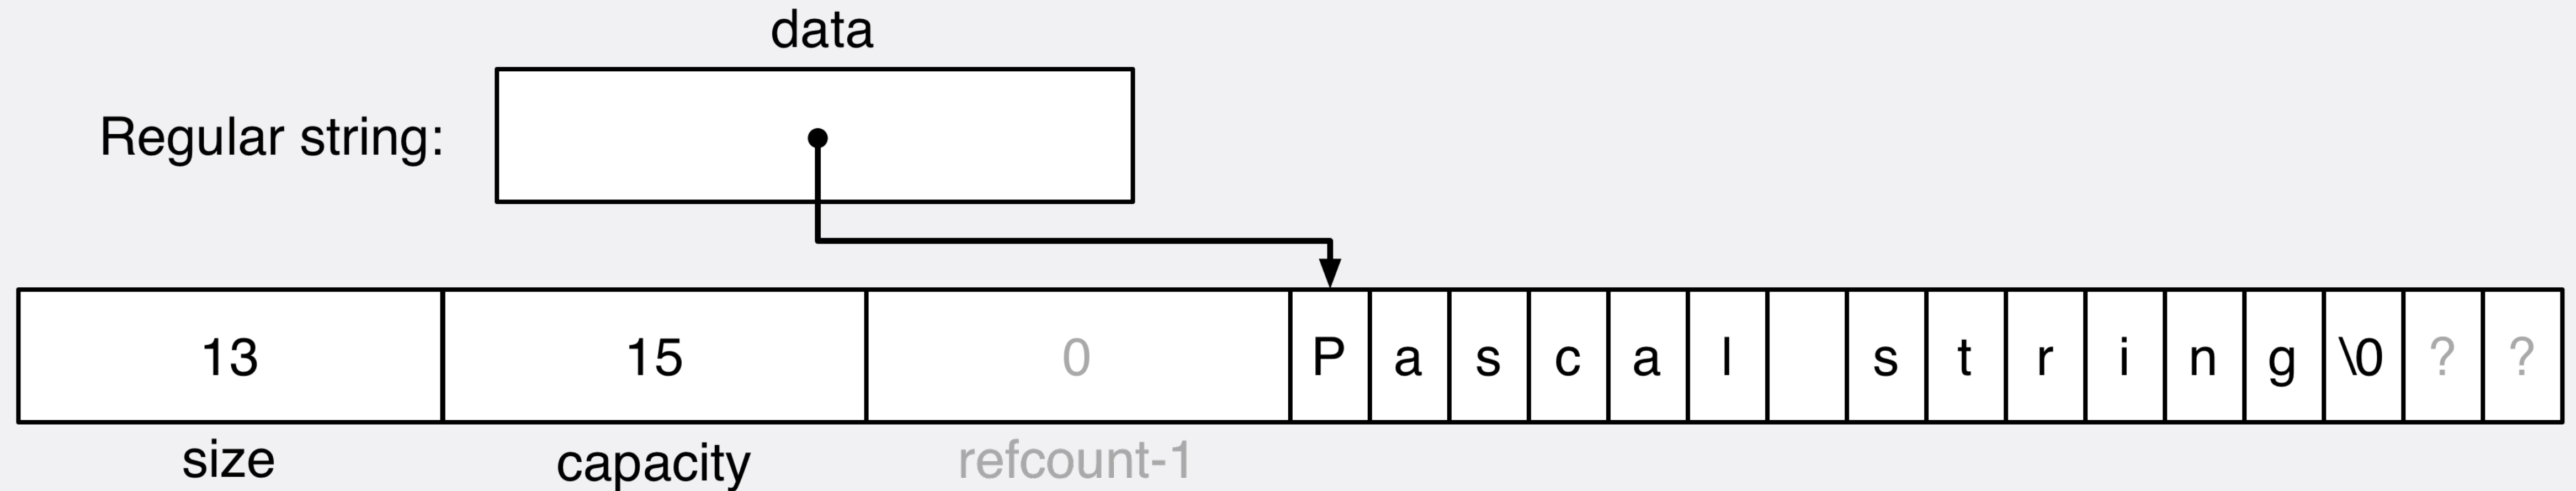
Why are strings important?

- `<string>` is the most-included file at Facebook
- Accounts for 18% of all CPU time spent in `std`
- There are simple ways to optimize strings

```
struct string {  
    int size;  
    int capacity;  
    char * data;  
};
```

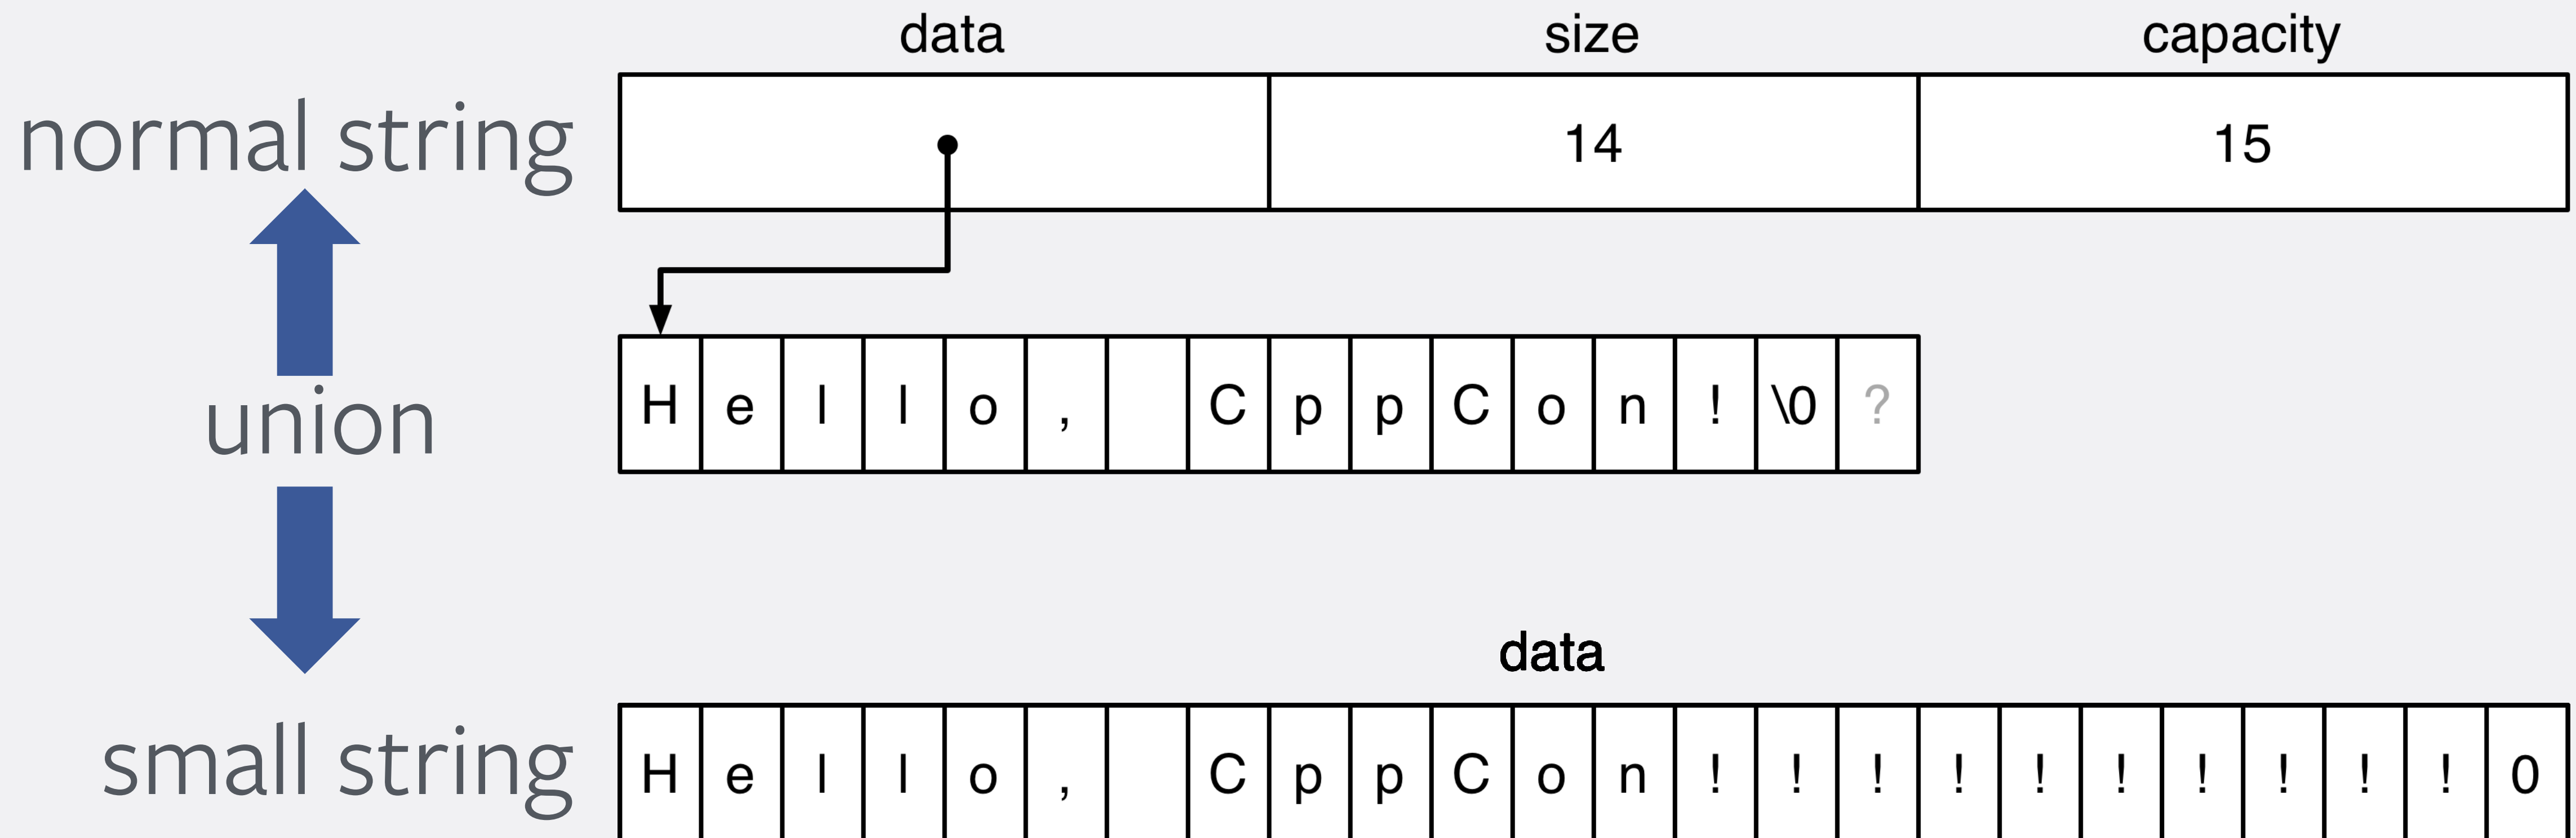


gcc string (version <5)



fbstring

@author Andrei Alexandrescu



Performance of fbstring

`gcc_string.size()`

`fbstring.size()`

```
movq    (%rdi), %rax
movq    -24(%rax), %rax
```

```
movabsq $-4611686018427387904, %rax
testq   %rax, 16(%rdi)
je      .L7
```

```
movq    8(%rdi), %rax
ret
.L7:
movsbq  23(%rdi), %rdx
movl    $23, %eax
subq    %rdx, %rax
ret
```

is_small

1.6ns

0.9ns

`std::string` replacement

- We replaced `std::string`'s implementation with `fbstring`'s
- `std::string` and `folly::fbstring` now have the same implementation, but are still different types

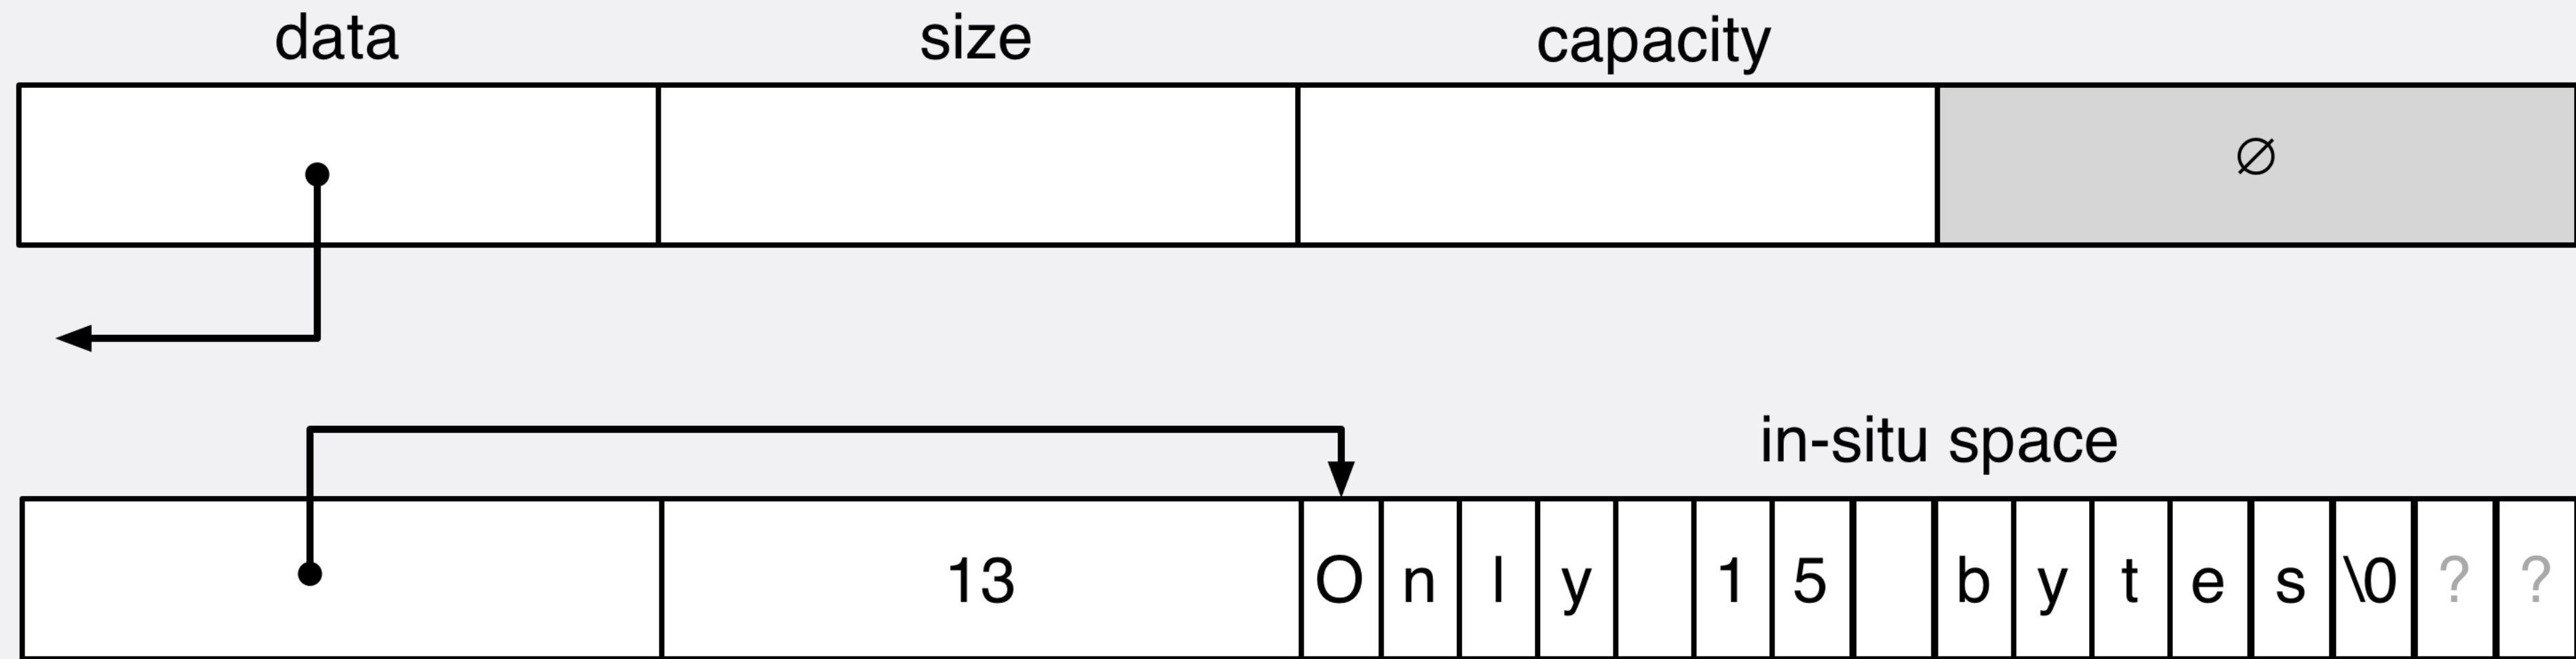
1% performance win

Killing the null terminator

- fbstring lazily wrote `\0`
- Added a mode that eagerly wrote `^` as terminator
- `c_str()`, `data()` will append `\0`

```
const Char * c_str() const {  
    ...  
    if (data[size()] != '\0')  
        data[size()] = '\0';  
    return data;  
}
```

gcc string (version >=5)



- + Has SSO
- + `data()`, `size()` very fast
- + Size, 32, is power of 2
- Only 15-byte capacity
- Move is no longer memcopy
- Size is 33% larger than fbstring

Strangeness no more!

- There are lots of different ways to implement strings
- Small-String-Optimization is in, Copy-On-Write is out
- Memory layout is important

facebook