

# What C++ Programmers Need to Know about Header `<random>`

Expanded edition, with minor corrections applied

WALTER E. BROWN, PH.D.

`< webrown.cpp @ gmail.com >`



Copyright © 2016 by Walter E. Brown. All rights reserved.

# Overview

---

- During this talk, we will:
  - Present the conceptual and technical underpinnings of the `<random>` header,
  - ✓ Demonstrate its most common correct usage pattern,
  - X Show the most common usage anti-pattern and explain what's wrong with it,
  - Explore some `<random>`-based toolkit designs,
  - Have a little `<random>` fun, and more!
- Some of this material is adapted from my WG21 papers:
  - N3551: “Random Number Generation in C++11” (2013).
  - N3847: “Random Number Generation is Not Simple!” (2014).
- Slides hidden (for lack of time) will be available online.

## A little about me

---

- B.A. (math's); M.S., Ph.D. (computer science).
- Professional programmer for ~~almost~~ 50 years, programming in C++ since 1982.
- Experienced in industry, academia, consulting, and research:
  - Founded a Computer Science Dept.; served as Professor and Dept. Head; taught and mentored at all levels.
  - Managed and mentored the programming staff for a reseller.
  - Lectured internationally as a software consultant and commercial trainer.
  - Retired from the Scientific Computing Division at Fermilab, specializing in C++ programming and in-house consulting.
- Not dead — still doing training & consulting. (Email me!)



## Emeritus participant in C++ standardization

- Written 100+ papers for WG21, proposing such now-standard C++ library features as `gcd/lcm`, `cbegin/cend`, and `common_type`, as well as the entirety of headers `<random>` and `<ratio>`.
  - Influenced such core language features as *alias templates*, *contextual conversions*, and *variable templates*; working on *requires-expressions*, operator synthesis, and more!
  - Conceived and served as Project Editor for ISO/IEC 29124 (Int'l Standard on Mathematical Special Functions in C++), now incorporated into C++17's `<cmath>`.
- Be forewarned: Based on my training and experience, I hold some rather strong opinions about computer software and programming methodology — these opinions are not shared by all programmers, but they should be! 😊



## Background Information re Randomness

Random numbers are a pain in the butt.  
However, they're terribly interesting,  
and very useful in many applications....

— Julianne Walker

Random numbers should not be generated  
with a method chosen at random.

— Donald Knuth

## In the beginning

- The first known algorithm is the *middle-square method*:
  - Maybe discovered by Brother Edvin, a Franciscan, circa 1245.
  - Discovered by John von Neumann (né Neumann János Lajos, 1903–1957) circa 1949.
  - Considered crude but fast, with obvious failure modes.
- The *Monte Carlo method* (a code name) of simulation:
  - Developed 1946+ by Stanislaw M. Ulam (1909–1984), coded on ENIAC by von Neumann, for the Manhattan Project.
  - Still heavily used today for modelling complicated processes via random numbers ...
  - Then statistically analyzing the resulting behavior.

## Uses of randomness today

- Monte Carlo and other simulations
- Genetic algorithms, randomized algorithms, Zobrist and other hashing algorithms
- Rearranging (permuting, shuffling):
  - testing sort algorithms
  - “randomizing” tournaments (sporting events/military draft/arranging candidates on voter ballots)
- Selection:
  - gaming (raffles/lotteries/slot machines)
  - calling on students fairly
  - aesthetics (art/music/poetry)
  - choosing performers on open-mike night
  - choosing quicksort’s pivot
- Statistical sampling:
  - mfg. quality control
  - drug screening
  - auditing
  - trials for new medical treatments/drugs
  - surveys
- Cryptography (has extra requirements not met by <random> )



“You have reached the Heisenberg Institute. Please hold and your call will be answered in random order.”



## Just recognizing randomness can be difficult

- Imagine an oracle gave you ten zeroes in a row. Random?
- “[T]here's nothing suspicious about ... a sequence of ten zeros. Ten values just isn't enough to draw any conclusions about the quality of a random number generator.”

— Pete Becker

- Expressed another way:



— Scott Adams

## Intuition fails most of us re randomness

- “... recent research [by P. Diaconis *et al.*] into coin flips has discovered that ... the chance of a coin landing in the same position as it started is about 51 percent. Heads facing up predicts heads; tails predicts tails.”  
— David E. Adler
- “I watched a guy at a Vegas casino flip a coin to tell if he was to ‘draw’ another card at a blackjack table.... The coin landed on its edge..... You could hear a pin drop!”  
— Joe Peach

## Achieving random behavior is astonishingly difficult

- “Statistical definitions [of randomness] involve the inability to predict outcomes or to find any pattern in a series of outcomes.”  
— Chris Wetzel
- “Randomness is really, really hard for computers.”  
— Jeff Atwood
- “[I]t's hard (and interesting) to get a computer to generate proper random numbers.”  
— Mads Haahr
- The randomness literature (including textbooks!) is littered with algorithms that ultimately proved to be inadequate at best, or outright wrong at worst.

## Training is sorely lacking

- Programming curricula rarely teach randomness.
- Random variates, widely thought to be simple, just aren't!
  - Example: if  $f()$  is uniformly distributed, so is  $2 * f()$ ;  
yet  $f() + f()$  is normally distributed, not uniformly.
  - So we must be careful with even simple refactorings!
- Programming puzzle (von Neumann, 1951):
  - Given a biased coin [unknown prob.  $0 < p < 1$  of heads] ...
  - How can we nonetheless algorithmically obtain fair results?
    - Solution in Wikipedia: “Fair coin” ...
    - But first think about it, okay? 😊

## Advice: tread extremely carefully!

- “Many generators have been written, most of them have demonstrably non-random characteristics, and some are embarrassingly bad.”  
— Stephen K. Park & Keith W. Miller
- “In the past, there have been some truly awful RNG algorithms published in well-known places, and then used by many people.”  
— Brad Lucier
- “The message ... is: *arbitrary operations on random numbers do not necessarily result in random output.*”  
— detly (blog)

And then we find ...

```
int getRandomNumber()  
{  
    return 4; // chosen by fair dice roll.  
              // guaranteed to be random.  
}
```

— Randall Munroe

(A classic case of  
“exactly what you asked for,  
but not what you wanted.”)

## What about rand( ) (in <stdlib.h>/<cstdlib>)?

- C11 says (in a footnote):
  - “There are no guarantees as to the quality of the random sequence produced and ...
  - “some implementations are known to produce sequences with distressingly non-random low-order bits. ...”
- C++17 (Committee Draft, N4604) says (in a Note) that:
  - “rand’s underlying algorithm is unspecified.
  - “Use of rand therefore continues to be nonportable, with unpredictable and oft-questionable quality and performance.”
- “[<random>] replaces the C random library the same way a computer replaces an abacus.” — Indi (blog)

# What's in Header <random>?

(Quite a Lot!)

Thanks to its uncompromising attention to generality and performance, one expert deemed [<random>] 'what every random number library wants to be when it grows up.'

— Bjarne Stroustrup




## <random> has something for everyone


- For casual users:
  - 1 preconfigured engine type
  - 2 configurable uniform distribution class templates
- For intermediate & advanced users:
  - 9 more preconfigured engine types
  - 18 more configurable distribution class templates
  - 1 URBG class for environmental sources of randomness
  - 1 class to help with engine seeding
- For experts:
  - 6 configurable engine/engine adapter class templates
  - 1 function template for authors of new distributions

## C++ nomenclature


- Traditional term *random number generator* is unfortunate:
  - It blurs the distinction between two kinds of functionality ...
  - So <random> uses more precise terminology.
- ① An engine is a means of obtaining a sequence of (ideally) unpredictable bits that are uniform (*i.e.*,  $p(0) = p(1)$ ).
- ② A distribution is a means of obtaining variates (values of desired aggregated shape) from such a sequence of bits.




Normal



Weibull



Poisson



Chi-square

## Engine types and objects

- An engine object  $e$  of type  $E$  gives  $n > 0$  bits per call  $e()$ :
  - Once initialized,  $e()$ 's output is deterministic, so is termed pseudo-random.
  - “Unpredictability is the ideal. Using a computer, we typically settle for very-very-very-very-hard-to-predict.”
- Some details:
  - Each call  $e()$  produces a *bit string* of nonzero length  $n$  ...
  - Encoded as a value of some unsigned  $E::\text{result\_type}$  ...
  - In the range  $[E::\text{min}() .. E::\text{max}()]$ , often  $[0 .. 2^n - 1]$ .
- But these details are rarely needed:
  - ☞ An engine object should be used nearly exclusively as a source of randomness (*e.g.*, for a distribution object).

## Distribution types and objects

- A distribution object produces a variate consistent with the distribution's result type and c'tor arguments:
  - *E.g.*, `std::uniform_int_distribution< long > d1{ low, high } ;`
  - *E.g.*, `std::normal_distribution< float > d2{ mean, stddev } ;`
- Examples:
  - ```
#include <random>
int roll_a_fair_die( ) {
    static std::default_random_engine      e{ } ;    // reproducible
    static std::uniform_int_distribution<int> d{ 1, 6 } ;
    return d(e);    // obtains a random variate
}
```
  - ```
constexpr std::size_t N = 1000;
int a[N];
std::generate( a + 0, a + N, roll_a_fair_die );
```

## <random> provides 10 pre-configured engine types

- Implementation-defined default\_random\_engine type:
  - Intended for “casual, inexperienced, and/or lightweight use.”
  - Vendor’s alg. choice (based on performance, size, quality, ...).
  - Needn’t give the same sequence across platforms.
- The other 9 alg’s are bit-for-bit portable across platforms:
  - Linear congruential engines: minstd\_rand0, minstd\_rand
  - Mersenne twister engines: mt19937, mt19937\_64
  - Subtract with carry engines: ranlux24\_base, ranlux48\_base
  - Discard block engines: ranlux24, ranlux48
  - Shuffle order engine: knuth\_b

## Why those 9 algorithms?

- Their characteristics (performance, size, quality, *etc.*) have been carefully studied for ~~many years~~ decades and are well described in standard references (*e.g.*, Knuth: *TAOCP* vol. 2):
  - *E.g.*, linear cong. is small and fast, but is of poorer quality ...
  - Than Mersenne twister, whose state is >600x larger.
- Underlying engine templates are for researchers, *etc.*:
  - Configurable via template parameters to provide additional engine types, *e.g.*, `linear_congruential_engine< ... >`.
  - But “Very few combinations of [template] parameters actually result in engines of decent quality.... The [non-expert] user who fools with [these] parameters ... should know that he is doing something foolish.” — Mark Fischler

## Engines can do more

- Some additional engine capabilities:
  - Engines are streamable (*e.g.*, to allow reproducibility or resuming an interrupted computation).
  - Engines of the same type can be compared for (in)equality.
  - Engines can be seeded (and also reseed( )ed) in more ways.
  - Engines can discard( ) (skip) generated values.
- But distributions use none of those extras:
  - Any type/object meeting a distribution's needs is (in C++17) termed a URBG (Uniform Random Bit Generator).
  - A URBG + all the extra capabilities is known as an engine.

## random\_device is a URBG

- Designed to be a standard interface to any available environmental/physical source of randomness:
  - *E.g.*, /dev/random or /dev/urandom.
  - *E.g.*, a device that samples atmospheric noise.
- Its c'tor argument is an implementation-defined string.
- Its entropy( ) member estimates the device's entropy, defined as:

$$S = - \sum_{i=0}^{n-1} P_i \log P_i$$



## Aside: entropy formula on Ludwig Boltzmann's tomb



— photo courtesy of Thomas Schneider

## Physical sources of randomness

- “The Ferranti Mark I computer, first installed in 1951, ...
  - “had a built-in instruction that put 20 random bits into the accumulator using a resistance noise generator; this feature had been recommended by A. M. Turing.” — Donald Knuth
- Hardware devices are commercially manufactured today:
  - Various hardware connections (USB, PCIe, ...) are available.
  - Recent CPUs support RDRAND, RDSEED instructions.
- Free online services:
  - [random.com](http://random.com), [randomnumbers.info](http://randomnumbers.info), [fourmilab.ch/hotbits](http://fourmilab.ch/hotbits)
  - Each seems amenable to a `random_device` interface.
  - But please take care, as there are also “psychic” services (offering lottery numbers and the like) for the unwary!

## <random> provides 20 distribution types in 5 families

- 2 Uniform distributions:

{uniform\_int,  
uniform\_real}\_distribution

- 4 Bernoulli distributions:

{bernoulli, binomial, geometric,  
negative\_binomial}\_distribution

- 5 Poisson distributions:

{poisson, exponential,  
gamma, weibull,  
extreme\_value}\_distribution

- 6 Normal distributions:

{normal, lognormal,  
chi\_squared, cauchy, fisher\_f,  
student\_t}\_distribution

- 3 Sampling distributions:

{discrete, piecewise\_constant,  
piecewise\_linear}\_distribution

? Why these dist's, not others?

Early rationale in N1588: "On  
Random-Number Distributions  
for C++0x" (2004).

## About std distributions

---

- Most of these are class templates whose template parameter is the desired type of the variate.
- Some deliver variates of only integer types, most others produce variates of only floating-point types:
  - *E.g.*, `discrete_distribution< integral_type >` .
  - *E.g.*, `uniform_real_distribution< floating_point_type >` .
  - Note: `bernoulli_distribution` gives only bool variates, so is a class (not a class template).
- While most std engines' results are bit-for-bit portable across platforms, distributions' results are not so specified:
  - Allows implementers to choose distribution algorithms (there are hundreds!) best for their target platforms.

## Important features: interoperability & extensibility

- By design, any URBG can be used with any distribution.
- The standard carefully defines URBG and dist. interfaces:
  - Thus, users can provide their own URBGs or distributions...
  - And theirs will seamlessly interoperate with the std ones.
- Users have taken advantage of this flexibility:
  - New engine algorithms are providing this interface, *e.g.*, PCG (M. O'Neill, 2015) and Random123 (J. Salmon, *et al.*, 2011).
  - And gcc ships with numerous contributed distributions:
    - k\_distribution
    - rice\_distribution
    - beta\_distribution
    - hoyt\_distribution
    - pareto\_distribution
    - arcsine\_distribution
    - logistic\_distribution
    - nakagami\_distribution
    - triangular\_distribution
    - von\_mises\_distribution
    - hypergeometric\_distribution
    - uniform\_on\_sphere\_distribution

## Writing your own distribution?

- <random> provides a utility for experts:
  - `template< class RealType, size_t bits, class URBG >`  
`RealType`  
`generate_canonical( URBG& g );`
  - Calls `g( )` as often as needed to get sufficient random bits, then maps those bits as uniformly as possible to the `RealType` range `[0 .. 1)`.
- Often “a useful step in the process of transforming a value generated by a uniform random bit generator into a value that can be delivered by a random number distribution.”

## How to test a new distribution?

- Apply statistical tests:
  - Check the sample mean and sample variance (if applicable).
  - Consider  $\chi^2$  and/or Kolmogorov-Smirnov tests.
  - Failure? Try again with a different engine or engine seed.
  - Remember that “Statistical tests are *supposed* to fail occasionally, just not often!” — John D. Cook
- ☺ Ensure that “If the numbers are not random, they are at least higgledy-piggledy.” — George Marsaglia

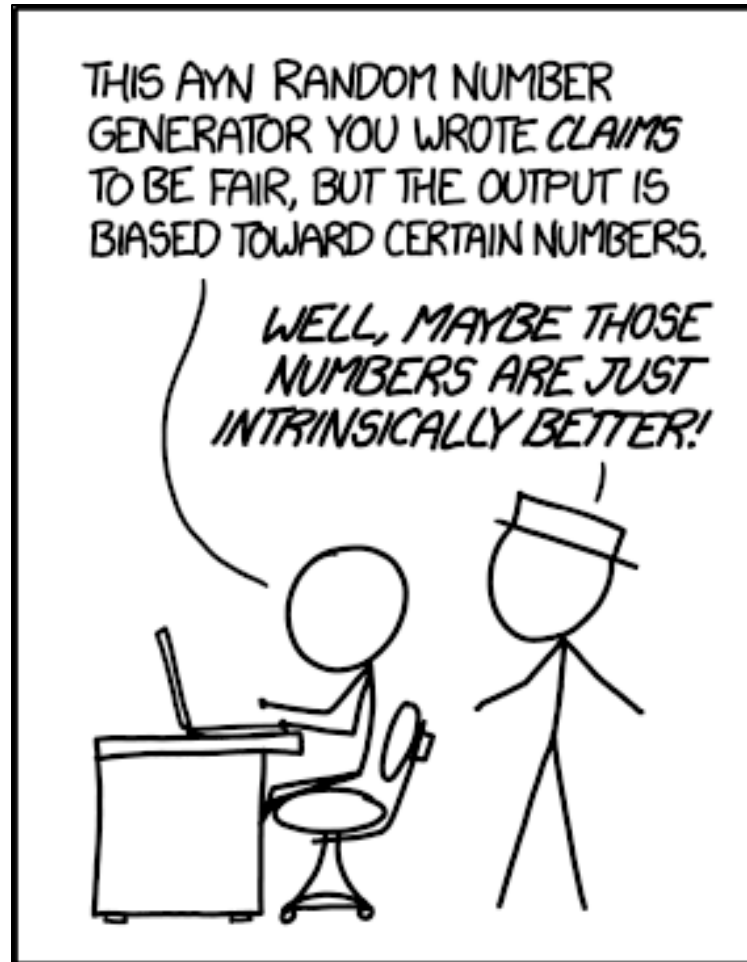
## Abusing the <random> Header

[Y]ou don't have to wait for someone to treat you bad[ly] repeatedly. All it takes is once, and if they get away with it that once ... it sets the pattern for the future.

— Jane Green



“Let not thy code introduce bias, lest ye die.”



— Randall Munroe

“When it comes to constructing [and using] random number generators, there is no such thing as getting the software *almost* right.”

— Stephen K. Park  
& Keith W. Miller

## Don't do this!

### ✗ Example of the most common anti-pattern:

- Let's mimic the roll of a fair die:

`e() % 6 + 1` // *seems innocuous; what's wrong?*

- ① Assumes that `e()` can produce at least 6 values (*i.e.*,  $2^n \geq 6$ ).
  - ② Since  $2^n$  isn't evenly divisible by 6, "this remainder operation makes lower numbers slightly more likely."
- "This is really awful!"
    - A floating-point (!) variation is "Hilariously non-uniform."
    - "It's the pigeonhole principle: if you have  $2^{32}$  pigeons, ... you can't put them into [6] pigeonholes without having more pigeons in some holes than others."

— Stephan T. Lavavej

## Such a roll is unfair (biased)

- “Anyone who [applies this remaindering technique] will be rewarded with a seemingly random sequence and be thrilled that their clever solution worked.
- “Unfo[r]tunately, this does not work ... because forcing the range in this way eliminates any chance of having a uniform distribution.
- “[T]o be correct, you must work with the distribution instead of destroy it.”

— Julianne Walker

- So let’s avoid such bigotry in programming, okay?

## Recommended Header <random> Usage

Wisdom is knowing what to do next, skill is knowing how to do it, and virtue is doing it.

— David Starr Jordan

## Let's do better

---

- Analogy: If we wanted a variate  $v$  in  $[1 .. 5]$ , we could roll a fair die, and re-roll (reject) if we get a 6:
  - `int v;`  
do `v = roll_the_die( ); while( v == 6 );`
  - This is a simple example of a *rejection algorithm*.
- More generally, if we wanted  $v$  in  $[low .. high]$ :
  - `int v;`  
do `v = e( ); while( v < low or v > high );` *// naïve rejection*
  - But what if  $e$ 's range doesn't include  $[low..high]$ ? And what's the performance if  $e$ 's range greatly exceeds  $[low..high]$ ? *Etc.*
- We're just re-inventing `std::uniform_int_distribution< > :`
  - Which is already part of `<random>!`

## Advice: when you want random variates

- Instantiate an engine and a distribution.
- For each variate, call your distribution object ...
- Passing it your engine object (as a source of randomness).
- Example: `auto variate = d(e);`
- Why? “It makes no sense to ask for a random number without some context. You need to specify random according to what distribution.” — John D. Cook

## Tip: manage engines and distributions as resources

- How many?
  - One of each is enough for many (most?) applications.
  - If so, can often treat them together:
    - *E.g.*, via pair, tuple, custom class, or class template.
    - *E.g.*, via bind, lambda, custom function, or function template.
- Engine objects are stateful and mutable, so:
  - Copy only when you want identical results in both places.
  - Usually pass by (non-const) reference, not by value.
  - Serialize access if using an engine in concurrent code:
    - Or consider `thread_local` engines (*i.e.*, 1 per thread).
- Ditto for distributions, so treat likewise:
  - *E.g.*, the *Box-Muller transform* calc's 2 normal variates at once.

## Initializing (seeding) an engine

- An engine's initial state is influenced by providing a seed (starting value) to its c'tor:
  - `std::default_random_engine e{ 13607 }; // still reproducible`
  - A seed's type matches (or converts to) `decltype( e( ) )`, a.k.a. `E::result_type` (the engine's unsigned result type).
- Engine's size is a factor:
  - Can't initialize Mersenne twister well with a single integer.
  - `std::seed_seq` takes one integer and attempts to make more entropy, but not very successfully; probably best avoided.



## Choosing good seeds can be challenging

- Reproducibility is critically important:
  - For some applications.
- But fatal for others:
  - “Players may be disappointed if your game always acts the same way given the same moves.” — John D. Cook
- Casting (to `E::result_type`) the current time or process id are common practices, but these are considered “low-quality seed data.”

What Else Can We Do with <random> ?

## Can sample( ) a range (population)

- New algorithm in C++17:
  - `template< class PopulationIter, class SamplerIter  
                    , class Size, class URBG >  
SampleIterator  
    sample( PopulationIter first, PopulationIter last  
            , SamplerIter out, Size wanted_size, URBG&& g );`
  - Uses one of two different algorithms, depending on iterators' capabilities.
- An `inplace_sample` algorithm is being proposed; will partition rather than copy.

## Can shuffle( ) a range (sequence)

- using card\_t = int;  
using deck\_t = array< card\_t, 52 >  
deck\_t deck; iota( begin(deck), end(deck), card\_t{ } );
- using engine\_t = default\_random\_engine;    *// vendor's choice*  
using seed\_t = engine\_t::result\_type;  
engine\_t e { seed\_t( time(0) ) };    *// a poor seed*  
shuffle( begin(deck), end(deck), e );    *// distribution is implicit*
- auto suit = [ ] ( card\_t c ) { return "♠♥♦♣" [ c / 13 ]; }  
auto rank = [ ] ( card\_t c ) { return "A23456789TJQK" [ c % 13 ]; }  
auto show = [ ] ( card\_t c ) { cout << ' ' << rank(c) << suit(c); }  
for\_each( begin(deck), end(deck), show );
- Note: the original STL random\_shuffle algorithms (which predate <random>) were deprecated in C++14 and now removed in C++17.

## Shuffling isn't always the right approach ☺

- ```
template< class RA, class Engine >  
void  
    shuffle_sort( RA b, RA e, Engine g )  
{  
    while( not is_sorted(b, e) )  
        shuffle( b, e, g );  
}
```

## Distribution parameters can vary per call, if desired

```
template< class RA, class URBG >
void shuffle( RA const b, RA const e, URBG&& g )
{
    using diff_t    = decltype( e - b );
    using dist_t    = uniform_int_distribution< diff_t >;
    using param_t = typename dist_t::param_type; // takes c'tor args

    static dist_t d{ }; // distribution's range supplied per call, below
    // invariant: items [b..b+m) are shuffled; items [b+m..b+L] are unshuffled
    diff_t const L = e - b - 1; // b[L] is the last item
    for( diff_t m{ }; m < L; ++m )
        iter_swap( b + m // the first unshuffled item
                   , b + d( g, param_t{m, L} ) // a random unshuffled item
                   );
}
```

## Consider an iterator interface to <random>

- using urbg\_t = default\_random\_engine;  
using variate\_t = double;  
using dist\_t = uniform\_real\_distribution<variate\_t>;
- urbg\_t g{ ... };  
dist\_t d{ }; *// default range is [ 0 .. 1 )*  
variate\_iterator<urbg\_t, dist\_t> it{ g, d }; *// see next page*
- *// create a random vector of extent N:*  
vector<variate\_t> v( N ); copy\_n( it, N, begin(v) );
- *// dot product with a random vector:*  
variate\_t prod{ inner\_product( begin(v), end(v), it, 0.0 ) };
- *// add noise to a vector:*  
transform( begin(v), end(v)  
          , begin(v)  
          , [&] ( variate\_t d ) { return d + \*it++; } );

## Sketch of an iterator interface to <random> 1

```
template< class URBG, class Dist >
class variate_iterator
: public iterator< input_iterator_tag, typename Dist::result_type >
{
private:
    using iter_t = variate_iterator;
    using val_t = typename Dist::result_type;
    using ptr_t = val_t const*;           // returned by op ->
    using ref_t = val_t const&;          // returned by op *

    URBG*      u{ nullptr }; // non-owning
    Dist*      d{ nullptr }; // non-owning
    val_t      v{ };         // latest variate
    bool       valid{ false };
```



## Sketch of an iterator interface to <random> 2

*// help:*

```
void step( ) noexcept { valid = false; }  
ref_t deref( ) { if( not valid ) v = (*d)(*u), valid = true;  
                return v; }
```

public:

*// construct:*

```
constexpr variate_iterator( ) noexcept = default;  
variate_iterator( URBG& u, Dist& d ) : u{ &u }, d{ &d } { ; }
```

*// dereference:*

```
ref_t operator* ( )      { return deref( ); }  
ptr_t operator->( )      { return &deref( ); }
```

*// advance:*

```
iter_t& operator++ ( )    { step( ); return *this; }  
iter_t operator++ ( int ) { iter_t t{*this}; step( ); return t; }  
}
```

# Final <random> Thoughts

Copyright © 2016 by Walter E. Brown. All rights reserved.

## *Caveat lector! 1*

- WG21 published an early spec. of `<random>` in TR1 (2007):
  - To get feedback from a wider community of users.
  - Worked! The adopted version is much improved over that early version in many important ways.
  - Many of these improvements rely on newer C++11 language features, or are otherwise not backward-compatible.
- But in the interim, early adopters commented on and offered advice regarding the then-new facility:
  - So the technical content of early writings (and, alas, of even some recent ones!) is often now at least somewhat suspect.
  - Even the historical content has often been wrong:
    - Boost.Random (J. Maurer) was our design's proof-of-concept!

## *Caveat lector! 2*

- We've also found many <random>-related writings:
  - That give programming advice based on what seems to be a questionable understanding of the meaning, behavior, and properties of randomness, and/or ...
  - That express opinion (commonly presented as fact) based on what seems to be a questionable understanding of <random>'s design.
- For all these reasons, I urge you:
  - Look carefully at such publications' sources and dates, ...
  - Before relying on their content.

## Some resources of interest

- M. Haahr, [random.org](http://random.org):
  - Brief tutorial re randomness.
  - Online environment-driven generation service.
- M. Klammler, P0205R0:
  - Recent proposal for enhanced (simplified) seeding; WG21 review under way.
- M. E. O'Neill, [pcg-random.org](http://pcg-random.org):
  - New PCG engine description, code, and formal theory paper.
  - Blogs describing other usability improvements, serving as a basis for recent proposal ...
  - P0347R0, by Song and O'Neill; WG21 review under way.



<random> is really elegantly designed. I love this header.  
— Stephan T. Lavavej

I think it's the best random number library design of all,  
by a mile. If I were a random number, I'd think I died  
and went to heaven.

— Andrei Alexandrescu

## But not everyone agrees

- “To the extent that anyone cares about C++11's random-number facility at all, the C++ community is polarized between two views....
  - “It's amazingly elegant, a shining example of separation of concerns [with] a pluggable and extensible architecture [that's] comprehensive and flexible.
  - “It's horrible to use, ... unnecessarily overengineered ..., completely unsuitable for beginners, and even seasoned programmers hate it.
- “Both camps are right.”

— M. E. O'Neill
- You'll have to make up your own mind. Thanks very much.

# What C++ Programmers Need to Know about Header `<random>`

## FIN

WALTER E. BROWN, PH.D.

`< webrown.cpp @ gmail.com >`