



U.S. DEPARTMENT OF
ENERGY



**UNIVERSITY OF
CALIFORNIA**



C++17 Parallel Algorithms

Bryce Adelstein Lebach


Computer Architecture Group, Computing Research Division

C++17 Parallel Algorithms

Bryce Adelstein Lebach

Computer Architecture Group, Computing Research Division

Desktop Compute Power (8-core 3.5GHz Sandy Bridge + AMD Radeon 6950)



© 2012 Adobe Systems Incorporated. All Rights Reserved.

18

Source: Sean Parent's C++Now 2012 Keynote



Hardware is increasingly parallel and increasingly diverse. Vendor-neutral parallel programming abstractions are desperately needed to avoid leaving performance on the table.

C++11/14 provide low-level concurrency primitives, but no real higher-level generic abstractions for parallel programming.

Q: What are standard algorithms?

The standard says the algorithms library
*describes components that C++
programs may use to perform
algorithmic operations on containers
and other sequences*

(25.1 [algorithms.general] p1)

Q: What are standard algorithms?

A: Generic operations on sequences.

Q: What are standard algorithms?

A: Generic operations on sequences.

(except `min`, `max`, etc)

Three types of standard algorithms:

- Non-modifying sequence operations.
- Mutating sequence operations.
- Sorting and related operations.

The Committee Draft of C++17 includes a new standard parallel algorithms library which provides parallelized versions of these sequence operations.

Q: What does parallel mean?

Bit-Level Parallelism

Instruction-Level Parallelism

Vector-Level Parallelism

Task-Level Parallelism

Process-Level Parallelism

implicit



explicit

Bit-Level Parallelism

Instruction-Level Parallelism

Vector-Level Parallelism

Task-Level Parallelism

Process-Level Parallelism

implicit



explicit

Bit-Level Parallelism

Instruction-Level Parallelism

Vector-Level Parallelism

Task-Level Parallelism

Process-Level Parallelism

implicit



explicit

Parallel algorithms library components:

- ExecutionPolicy concept.
- Three standard execution policies.
- New ExecutionPolicy overloads for most existing standard algorithms.
- New “unordered” algorithms based on existing “ordered” algorithms.
- Fused algorithms.

An `ExecutionPolicy` describes how a generic algorithm may be parallelized.

They allow programmers to request parallelism and describe constraints.

Standard execution policies:

- `std::execution::par` – operations are indeterminately sequenced in the calling thread.
- `std::execution::par` – operations are indeterminately sequenced with respect to each other within the same thread.
- `std::execution::par_unseq` – operations are unsequenced with respect to each other and possibly interleaved.

Suppose we are using the following binary operation with `std::transform`:

```
double multiply(double x, double y)
{ return x * y; }
```

```
std::transform(
    // "Left" input sequence.
    x.begin(), x.end(),
    y.begin(), // "Right" input sequence.
    x.begin(), // Output sequence.
    multiply);
```

Suppose we are using the following binary operation with `std::transform`:

```
load x[i] to a scalar register  
load y[i] to a scalar register  
multiply x[i] and y[i]  
store the result to x[i]
```

std::par

load $x[i]$ to a scalar register
load $y[i]$ to a scalar register
multiply $x[i]$ and $y[i]$
store the result to $x[i]$

load $x[i+1]$ to a scalar register
load $y[i+1]$ to a scalar register
multiply $x[i+1]$ and $y[i+1]$
store the result to $x[i+1]$

load $x[i+2]$ to a scalar register
load $y[i+2]$ to a scalar register
multiply $x[i+2]$ and $y[i+2]$
store the result to $x[i+2]$

load $x[i+3]$ to a scalar register
load $y[i+3]$ to a scalar register
multiply $x[i+3]$ and $y[i+3]$
store the result to $x[i+3]$

std::par

```
load x[i ] to a scalar register
load y[i ] to a scalar register
multiply x[i ] and y[i ]
store the result to x[i ]
load x[i+1] to a scalar register
load y[i+1] to a scalar register
multiply x[i+1] and y[i+1]
store the result to x[i+1]
load x[i+2] to a scalar register
load y[i+2] to a scalar register
multiply x[i+2] and y[i+2]
store the result to x[i+2]
load x[i+3] to a scalar register
load y[i+3] to a scalar register
multiply x[i+3] and y[i+3]
store the result to x[i+3]
```

std::par_unseq

```
load x[i ] to a scalar register
load x[i+1] to a scalar register
load x[i+2] to a scalar register
load x[i+3] to a scalar register
load y[i ] to a scalar register
load y[i+1] to a scalar register
load y[i+2] to a scalar register
load y[i+3] to a scalar register
multiply x[i ] and y[i ]
multiply x[i+1] and y[i+1]
multiply x[i+2] and y[i+2]
multiply x[i+3] and y[i+3]
store the result to x[i ]
store the result to x[i+1]
store the result to x[i+2]
store the result to x[i+3]
```

std::par

```
load x[i ] to a scalar register
load y[i ] to a scalar register
multiply x[i ] and y[i ]
store the result to x[i ]
load x[i+1] to a scalar register
load y[i+1] to a scalar register
multiply x[i+1] and y[i+1]
store the result to x[i+1]
load x[i+2] to a scalar register
load y[i+2] to a scalar register
multiply x[i+2] and y[i+2]
store the result to x[i+2]
load x[i+3] to a scalar register
load y[i+3] to a scalar register
multiply x[i+3] and y[i+3]
store the result to x[i+3]
```

std::par_unseq

```
load x[i:i+3] to a vector register
load y[i:i+3] to a vector register
multiply x[i:i+3] and y[i:i+3]
store the results to x[i:i+3]
```

New ExecutionPolicy overloads for most existing standard algorithms.

adjacent_difference	is_sorted[_until]	rotate[_copy]
adjacent_find	lexicographical_compare	Search[_n]
all_of	max_element	set_difference
any_of	merge	set_intersection
copy[_if _n]	min_element	set_symmetric_difference
count[_if]	minmax_element	set_union
equal	mismatch	sort
fill[_n]	move	stable_partition
find[_end _first_of _if _if_not]	none_of	stable_sort
for_each	nth_element	swap_ranges
generate[_n]	partial_sort[_copy]	transform
includes	partition[_copy]	uninitialized_copy[_n]
inplace_merge	remove[_copy _copy_if _if]	uninitialized_fill[_n]
is_heap[_until]	replace[_copy _copy_if _if]	unique
is_partitioned	reverse[_copy]	unique_copy


```
std::vector<T> x = // ...
```

```
std::sort(std::execution::par,  
          x.begin(), x.end());
```

```
std::vector<T> x = // ...
```

```
std::for_each(std::execution::par_unseq,  
              x.begin(), x.end(), process);
```

```
std::vector<T> x = // ...
```

```
#pragma omp parallel for simd  
for (std::size_t i = 0; i < x.size(); ++i)  
    process(x[i]);
```

New “unordered” algorithms based on existing “ordered” algorithms.

- `std::reduce`
- `std::inclusive_scan`
- `std::exclusive_scan`
- `std::transform_reduce`

`std::reduce - unordered` `std::accumulate`

```
T v = std::reduce([ep,  
                  first, last,  
                  [init,] [op])
```

```
std::accumulate:  
  first, acc = init  
  then for every it in [first, last) in order  
  acc = binary_op(acc, *it)
```

```
std::reduce:  
  GSUM(binary_op, init, *first, ...)
```

Commutativity: Changing the order of operations does not change the result.

- Integer Addition: $x + y == y + x$
- Integer Multiplication: $xy == yx$
- Integer Subtraction: $x - y \neq y - x$

Associativity: The grouping of operations does not change the result.

- Integer Addition: $(x + y) + z == x + (y + z)$
- Integer Multiplication: $(xy)z == x(yz)$
- Integer Subtraction: $(x - y) - z == x - (y - z)$

$$\text{GNSUM}(\text{op}, a^1, \dots, a^N) = \begin{cases} a^1 & N == 1 \\ \text{op}(\text{GNSUM}(\text{op}, a^1, \dots, a^k), & \text{otherwise} \\ \quad \text{GNSUM}(\text{op}, a^{k+1}, \dots, a^N)) \end{cases}$$

$$\text{GSUM}(\text{op}, a^1, \dots, a^N) == \text{GNSUM}(\text{op}, b^1, \dots, b^N)$$

where b^1, \dots, b^N may be any permutation of b^1, \dots, b^N

```
std::vector<double> x{1e-2, 1e-1, 1e0, 1e1, 1e2};  
  
// sum ~= 111.111  
double sum = std::accumulate(x.begin(), x.end(), 0.0);
```

```
std::vector<double> x{1e-2, 1e-1, 1e0, 1e1, 1e2};
```

```
// sum ~= 111.111
```

```
double sum = 0.0;
```

```
    sum = sum + x[0];
```

```
    sum = sum + x[1];
```

```
    sum = sum + x[2];
```

```
    sum = sum + x[3];
```

```
    sum = sum + x[4];
```

```
std::vector<double> x{1e-2, 1e-1, 1e0, 1e1, 1e2};  
  
// sum ~= 111.111  
double sum = std::reduce(x.begin(), x.end(), 0.0);
```

```
std::vector<double> x{1e-2, 1e-1, 1e0, 1e1, 1e2};
```

```
// sum ~= 111.111
```

```
double sum = 0.0;
```

```
    sum = sum + x[0];
```

```
    sum = sum + x[1];
```

```
    sum = sum + x[2];
```

```
    sum = sum + x[3];
```

```
    sum = sum + x[4];
```

```
std::vector<double> x{1e-2, 1e-1, 1e0, 1e1, 1e2};
```

```
// sum ~= 111.111
```

```
double sum = 0.0;
```

```
    sum = sum + x[1];
```

```
    sum = sum + x[0];
```

```
    sum = sum + x[2];
```

```
    sum = sum + x[4];
```


```
    sum = sum + x[3];
```

```
std::vector<double> x{1e-2, 1e-1, 1e0, 1e1, 1e2};
```

```
// GNSUM(plus, x[2], x[3], x[4])  
double sum0 = x[2] + x[3] + x[4];
```

```
// GNSUM(plus, 0.0, x[0], x[1])  
double sum1 = 0.0 + x[0] + x[1];
```

```
// GNSUM(plus, 0.0, x[0], ..., x[4])  
// = plus(GNSUM(plus, x[2], x[3], x[4]),  
//        GNSUM(plus, 0.0, x[0], x[1]))  
double sum = sum0 + sum1;
```




```
std::vector<double> x{1e-2, 1e-1, 1e0, 1e1, 1e2};
```

```
// GNSUM(plus, x[0], x[4], x[2])  
double sum0 = x[0] + x[4] + x[2];
```

```
// GNSUM(plus, 0.0, x[3], x[1])  
double sum1 = 0.0 + x[3] + x[1];
```

```
// GNSUM(plus, 0.0, x[0], ..., x[4])  
// = plus(GNSUM(plus, x[0], x[4], x[2]),  
//        GNSUM(plus, 0.0, x[3], x[1]))  
double sum = sum0 + sum1;
```



`std::inclusive_scan - unordered std::partial_sum.`

```
OutputIt it =  
    std::inclusive_scan([ep,]  
                        first, last, output,  
                        [op,] [init]);
```

`std::inclusive_scan - unordered std::partial_sum.`

```
OutputIt it =  
    std::inclusive_scan([ep,]  
                        first, last, output,  
                        [op,] [init]);
```

```
*(output)      = *first;  
*(output+1)    = *first + *(first+1);  
*(output+2)    = *first + *(first+1) + *(first+2);  
// ...
```

`std::exclusive_scan()` - exclusive prefix sum.

```
OutputIt it =  
    std::exclusive_scan([ep,  
                        first, last, output,  
                        [op,] [init]]);
```

```
*(output)      = init;  
*(output+1)    = init + *first;  
*(output+2)    = init + *first + *(first+1);  
// ...
```

Fused algorithms:

- `transform_reduce`
- `transform_inclusive_scan`
- `transform_exclusive_scan`

C++17 Parallel Algorithms

Bryce Adelstein Lebach

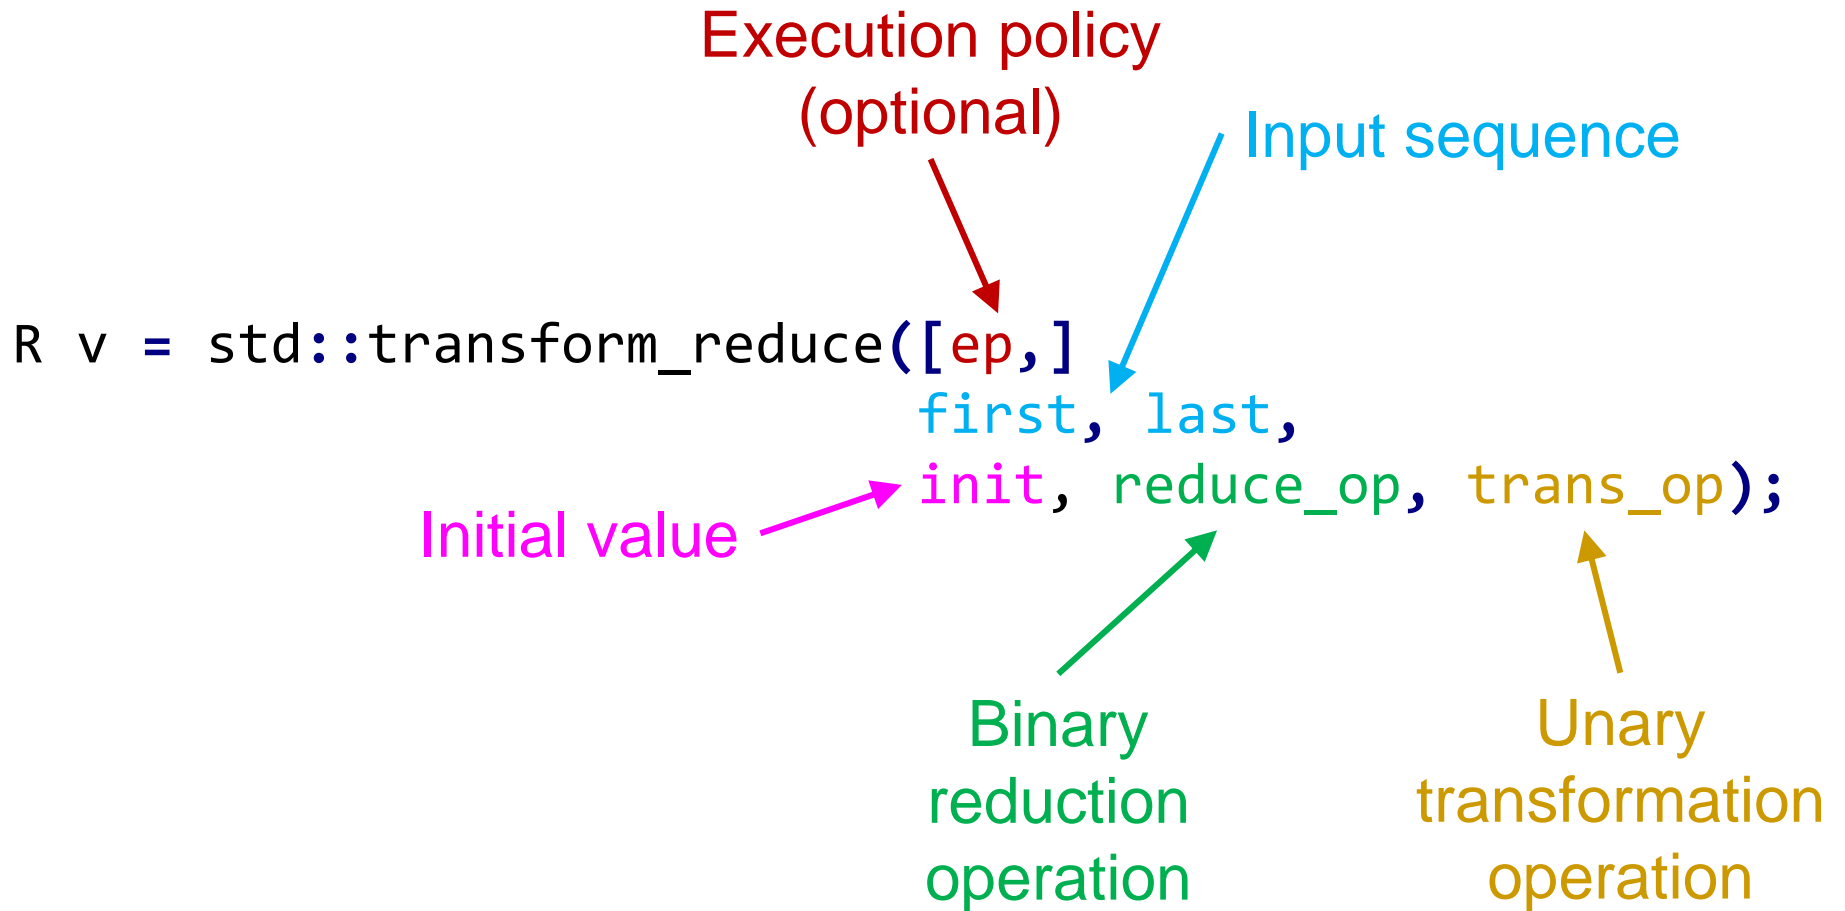
Computer Architecture Group, Computing Research Division

~~C++17 Parallel Algorithms~~

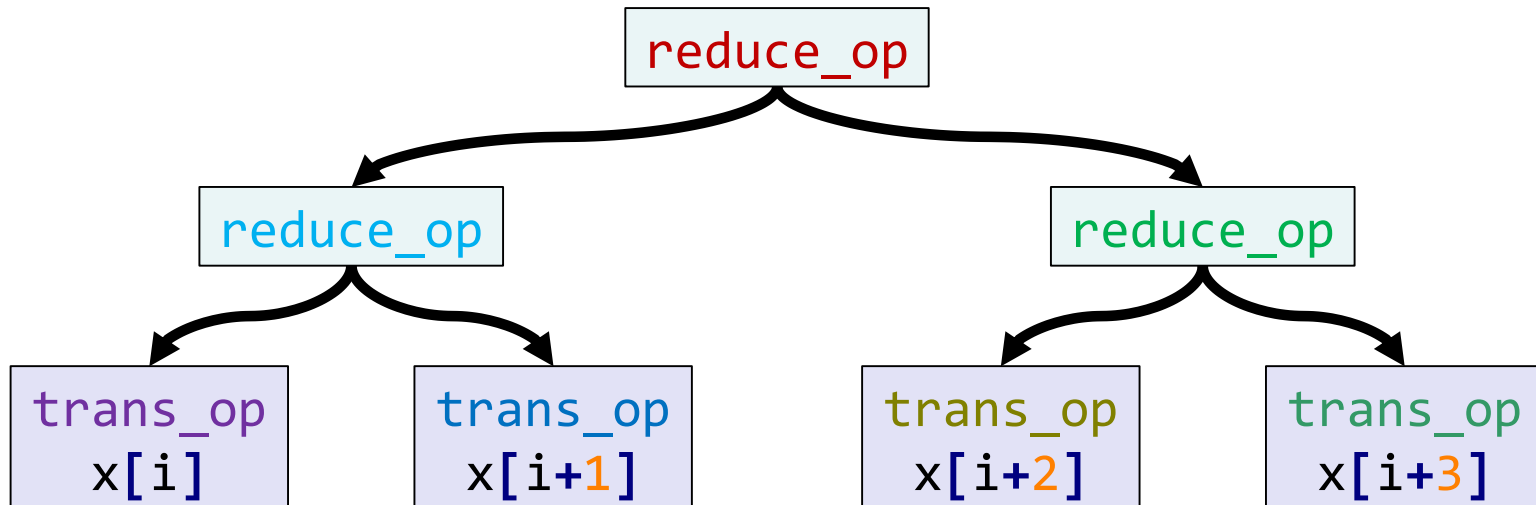
Bryce Adelstein LeBach

Computer Architecture Group, Computing Research Division


std::transform_reduce



```
R v = reduce_op(  
  reduce_op(trans_op(x[i]), trans_op(x[i+1])),  
  reduce_op(trans_op(x[i+2]), trans_op(x[i+3])));
```



$$R \ v = (x[i]^2 + x[i+1]^2) + (x[i+2]^2 + x[i+3]^2) \dots$$



```
std::vector<double> x = // ...
```

```
double norm =  
    std::sqrt((x[0] * x[0]) + (x[1] * x[1]) + /* ... */);
```

```

std::vector<double> x = // ...

double norm =
    std::sqrt(
        std::transform_reduce(
            std::execution::par_unseq,

            // Input sequence.
            x.begin(), x.end(),

            // Initial reduction value.
            double(0.0),

            // Binary reduction op.
            [] (double x1, double xr) { return x1 + xr; },

            // Unary transform op.
            [] (double x) { return x * x; }
        )
    );

```

Execution policy (optional)

Input sequence #1

Input sequence #2

```
R v = std::transform_reduce([ep,] first1, last2, first2, init, reduce_op, trans_op);
```

Initial value

Binary reduction operation

Binary transformation operation

```
std::vector<double> x = // ...  
std::vector<double> y = // ...  
  
double dot_product =  
    (x[0] * y[0]) + (x[1] * y[1]) + // ...
```

```
std::vector<double> x = // ...  
std::vector<double> y = // ...
```

```
double dot_product = std::transform_reduce(  
    std::execution::par_unseq, x.begin(), x.end(), y.begin());
```



```
std::size_t word_count(std::string_view s) {  
    // Goal: Count the number of word "beginnings" in the  
    // input sequence.  
  
    // ...
```

```
bool is_word_beginning(char left, char right) {  
    // If left is a space and right is not, we've hit a  
    // new word.  
    return std::isspace(left) && !std::isspace(right);  
}
```

```
std::size_t word_count(std::string_view s) {  
    if (s.empty()) return 0;  
  
    // ...  
}
```

```
std::size_t word_count(std::string_view s) {  
    if (s.empty()) return 0;  
  
    // If the first character is not a space, then it's the  
    // beginning of a word.  
    std::size_t wc = (!std::isspace(s.front()) ? 1 : 0);  
  
    // ...  
}
```

```

std::size_t word_count(std::string_view s) {
    // ...

    // Count the number of characters that start a new word
    WC +=
        std::transform_reduce(
            std::execution::par_unseq,

            // "Left" input: s[0], s[1], ..., s[s.size() - 2]
            s.begin(), s.end() - 1,
            // "Right" input: s[1], s[2], ..., s[s.size() - 1]
            s.begin() + 1,

            // ...

```

```

std::size_t word_count(std::string_view s) {
    // ...

    // Count the number of characters that start a new word
    WC +=
        std::transform_reduce(
            std::execution::par_unseq,

            // "Left" input: s[0], s[1], ..., s[s.size() - 2]
            s.begin(), s.end() - 1,
            // "Right" input: s[1], s[2], ..., s[s.size() - 1]
            s.begin() + 1,

            std::size_t(0), // Initial value for reduction.

            std::plus<std::size_t>(), // Binary reduction op.
            is_word_beginning         // Binary transform op: Return
                                    // 1 when we hit a new word.
        );

    // ...

```

Input sequence:

"Whose woods these are I think I know.\n"

"His house is in the village though; \n"

"He will not see me stopping here \n"

"To watch his woods fill up with snow.\n"

First Stanza of Stopping by Woods on a Snowy Evening, Robert Frost

Post-transform pseudo-sequence:

0000010000010000010001010000010100000
10001000001001001000100000000100000000
10010000100010001001000000000100000000
10010000010001000001000010010000100000

Input sequence:

"Whose woods these are I think I know.\n"

"His house is in the village though; \n"

"He will not see me stopping here \n"

"To watch his woods fill up with snow.\n"

First Stanza of Stopping by Woods on a Snowy Evening, Robert Frost

Post-transform pseudo-sequence:

1 + 1 + 1 + 1+1 + 1+1 +
1 + 1 + 1 +1 +1 + 1 + 1 +
1 +1 + 1 + 1 + 1 +1 + 1 +
1 +1 + 1 + 1 + 1 + 1 +1 + 1


```

bool is_word_beginning(char left, char right) {
    return std::isspace(left) && !std::isspace(right);
}

std::size_t word_count(std::string_view s) {
    if (s.empty()) return 0;

    std::size_t wc = (!std::isspace(s.front()) ? 1 : 0);

    wc +=
        std::transform_reduce(
            std::execution::par_unseq,
            s.begin(), s.end() - 1,
            s.begin() + 1,
            std::size_t(0),
            std::plus<std::size_t>(),
            is_word_beginning
        );

    return wc;
}

```

```

bool is_word_beginning(char left, char right) {
    return std::isspace(left) && !std::isspace(right);
}

std::size_t word_count(std::string_view s) {
    if (s.empty()) return 0;

    std::size_t wc =
        std::transform_reduce(
            std::execution::par_unseq,
            s.begin(), s.end() - 1,
            s.begin() + 1,
            std::size_t(!std::isspace(s.front()) ? 1 : 0),
            std::plus<std::size_t>(),
            is_word_beginning
        );

    return wc;
}

```

Sparse histogram:

- Goal: Find all the unique values in a sequence and count the number of times they occur.
- Example:
 - Input: a, b, c, c, a, a, b, b, b, b, e
 - Output keys: [a, b, c, e]
 - Output counts: [2, 5, 3, 1]

```
auto sparse_histogram(std::vector<T> const& x) {  
    std::vector<T>          hist_keys;  
    std::vector<std::size_t> hist_counts;  
  
    // ...  
}
```

```
auto sparse_histogram(std::vector<T> const& x) {  
    std::vector<T>          hist_keys;  
    std::vector<std::size_t> hist_counts;  
  
    if (x.empty())  
        return std::make_tuple(std::move(hist_keys),  
                                std::move(hist_counts));  
  
    // ...  
}
```

```
auto sparse_histogram(std::vector<T> const& x) {  
    std::vector<T>          hist_keys;  
    std::vector<std::size_t> hist_counts;  
  
    if (x.empty())  
        return std::make_tuple(std::move(hist_keys),  
                                std::move(hist_counts));  
  
    // Sort x to bring equal elements together.  
    std::sort(std::execution::par_unseq, x.begin(), x.end());  
  
    // ...  
}
```

```
auto sparse_histogram(std::vector<T> const& x) {  
    // ...  
  
    // Count the number of unique elements.  
    std::size_t num_unique_elements = // ...  
  
    // ...
```

```

auto sparse_histogram(std::vector<T> const& x) {
    // ...

    // Count the number of unique elements.
    std::size_t num_unique_elements = std::transform_reduce(
        std::execution::par_unseq,

        x.begin(), x.end() - 1, // x[0], x[1], ..., x[x.size() - 2]
        x.begin() + 1,          // x[1], x[2], ..., x[x.size() - 1]

        std::size_t(1), // x isn't empty, so 1 unique key minimum.
        std::plus<std::size_t>() // Reduction operation.

        // Transform op: Return 1 if right is a new unique element.
        [] (auto&& left, auto&& right)
        // If the right is not equal to the left, then we've
        // hit the next unique element.
        { return left != right; },
    );

    // ...
}

```



```
auto sparse_histogram(std::vector<T> const& x) {  
    // ...  
  
    // Allocate storage.  
    hist_keys.resize (num_unique_elements);  
    hist_counts.resize(num_unique_elements);  
  
    // ...  
}
```

```

auto sparse_histogram(std::vector<T> const& x) {
    // ...

    // Count the number of occurrences of each unique key.
    hpx::reduce_by_key(
        hpx::execution::par_unseq,

        // Input key sequence.
        x.begin(), x.end(),
        // Input value sequence.
        boost::constant_iterator<std::size_t>(1),

        // Output key sequence.
        hist_keys.begin(),
        // Output value sequence.
        hist_counts.begin()
    );

    // ...
}

```

```

auto sparse_histogram(std::vector<T> const& x) {
    // ...

    // Count the number of occurrences of each unique key.
    hpx::reduce_by_key(
        hpx::execution::par_unseq,

        // Input key sequence.
        x.begin(), x.end(),
        // Input value sequence.
        boost::constant_iterator<std::size_t>(1),

        // Output key sequence.
        hist_keys.begin(),
        // Output value sequence.
        hist_counts.begin()
    );

    // ...
}

```

```

auto sparse_histogram(std::vector<T> const& x) {
    // ...

    // Count the number of occurrences of each unique key.
    hpx::reduce_by_key(
        hpx::execution::par_unseq,

        // Input key sequence.
        x.begin(), x.end(),
        // Input value sequence.
        boost::constant_iterator<std::size_t>(1),

        // Output key sequence.
        hist_keys.begin(),
        // Output value sequence.
        hist_counts.begin()
    );

    // ...
}

```

```

auto sparse_histogram(std::vector<T> const& x) {
    // ...

    hpx::reduce_by_key(
        hpx::execution::par_unseq,

        // Input key sequence.
        x.begin(), x.end(),
        // Input value sequence.
        boost::constant_iterator<std::size_t>(1),

        // Output key sequence.
        hist_keys.begin(),
        // Output value sequence.
        hist_counts.begin()
    );

    return std::make_tuple(std::move(hist_keys),
                           std::move(hist_counts));
}

```

Parallel algorithms exception handling

- If an element access function exits via an uncaught exception, `std::terminate` is called.
- Parallel algorithms may also throw `std::bad_alloc` if temporary memory resources are needed for execution and none are available.

Acknowledgments

HPX: github.com/STELLAR-GROUP/hpx

Thrust: github.com/thrust/thrust

Boost.Compute: github.com/boostorg/compute

Jared Hoberock

Michael Garland

Grant Mercer

Hartmut Kaiser

John Biddiscombe